

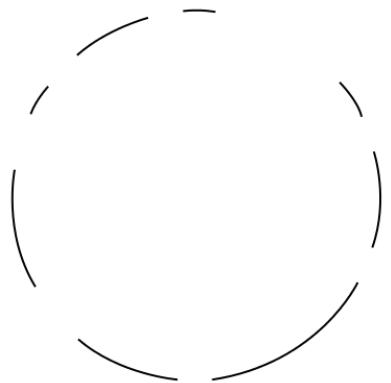
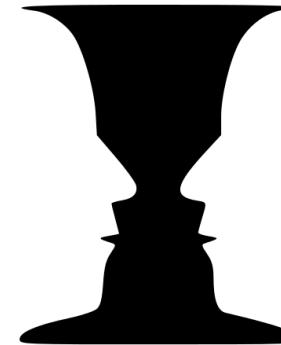
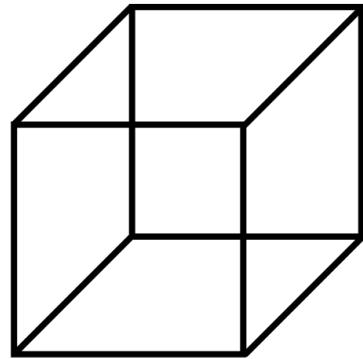
# **합성곱 신경망**

# **Convolutional Neural Networks**

**오영민,**  
**Ph.D.**

# Problem of Perception 지각의 문제

지각(perception)은 감각 정보에 대한 능동적인 해석, 혹은 번역이다.

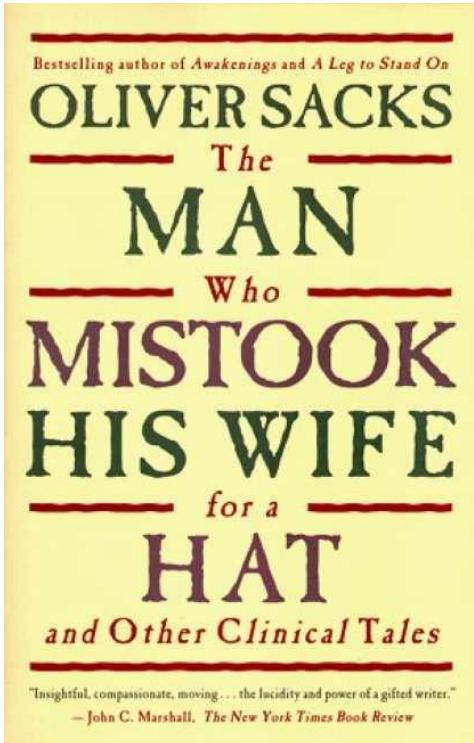


<https://en.wikipedia.org/wiki/Perception>

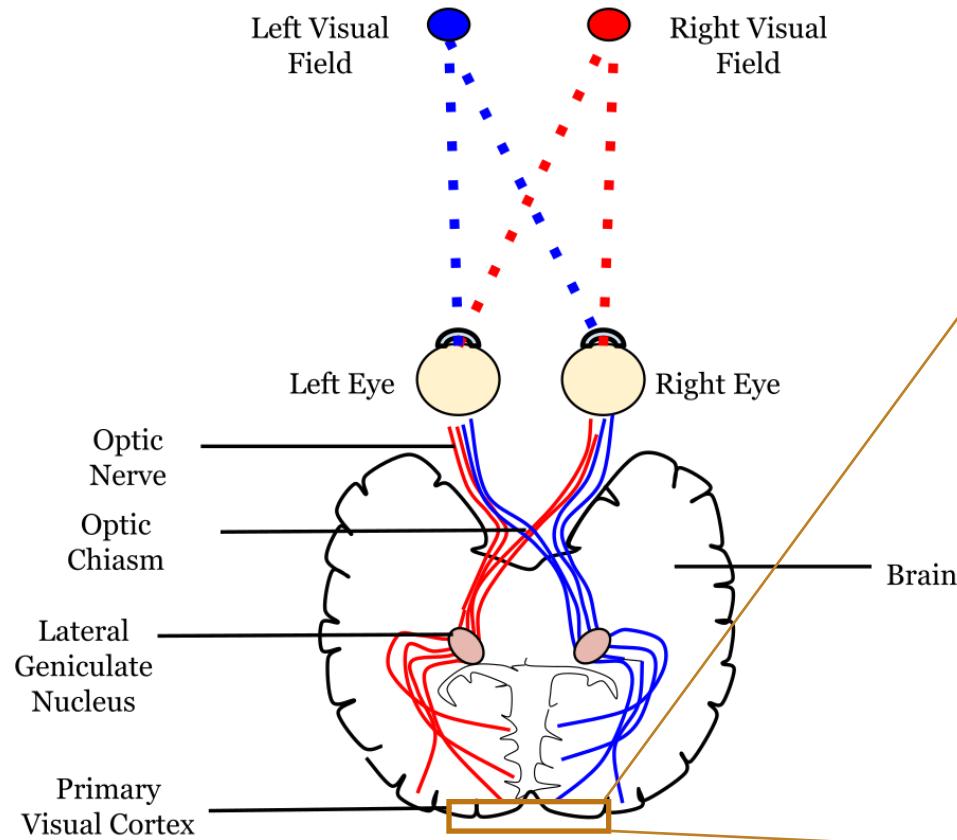
# Visual Agnosia 시각실인증

시각 기능과 인지 능력이 온전함에도 불구하고, 특정 사물을 인식하지 못하는 증상.  
시각적 인식(지각)이 sub-conscious level에서 이루어지는 고도의 기능임을 알 수 있다.  
예) 안면실인증

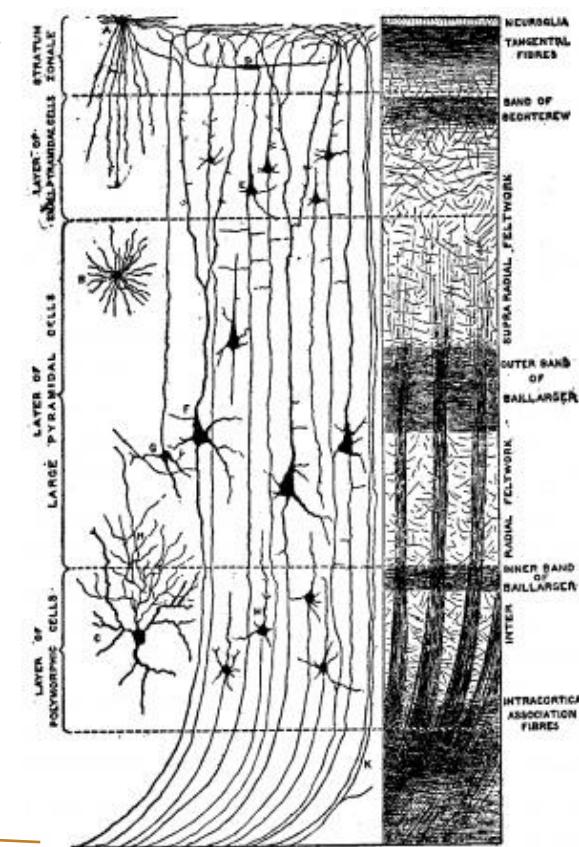
<https://mozanunal.com/2019/11/img2sh/>



# Visual Recognition of Brain 뇌의 시각 정보 인식



[https://upload.wikimedia.org/wikipedia/commons/4/4d/Neural\\_pathway\\_diagram.svg](https://upload.wikimedia.org/wikipedia/commons/4/4d/Neural_pathway_diagram.svg)



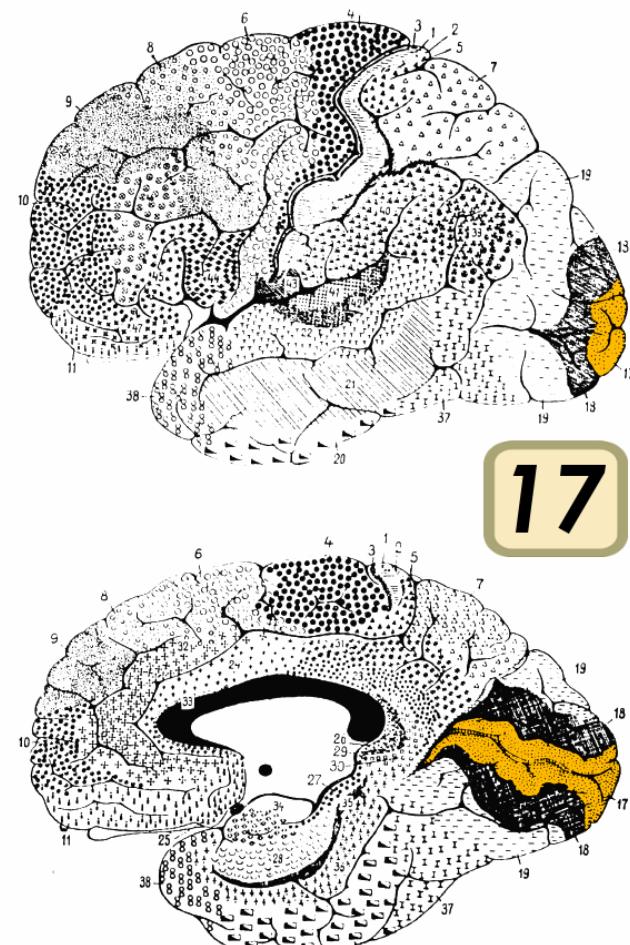
[https://en.wikipedia.org/wiki/Cerebral\\_cortex](https://en.wikipedia.org/wiki/Cerebral_cortex)

# Layers of Cerebral Cortex 대뇌피질의 층 구조

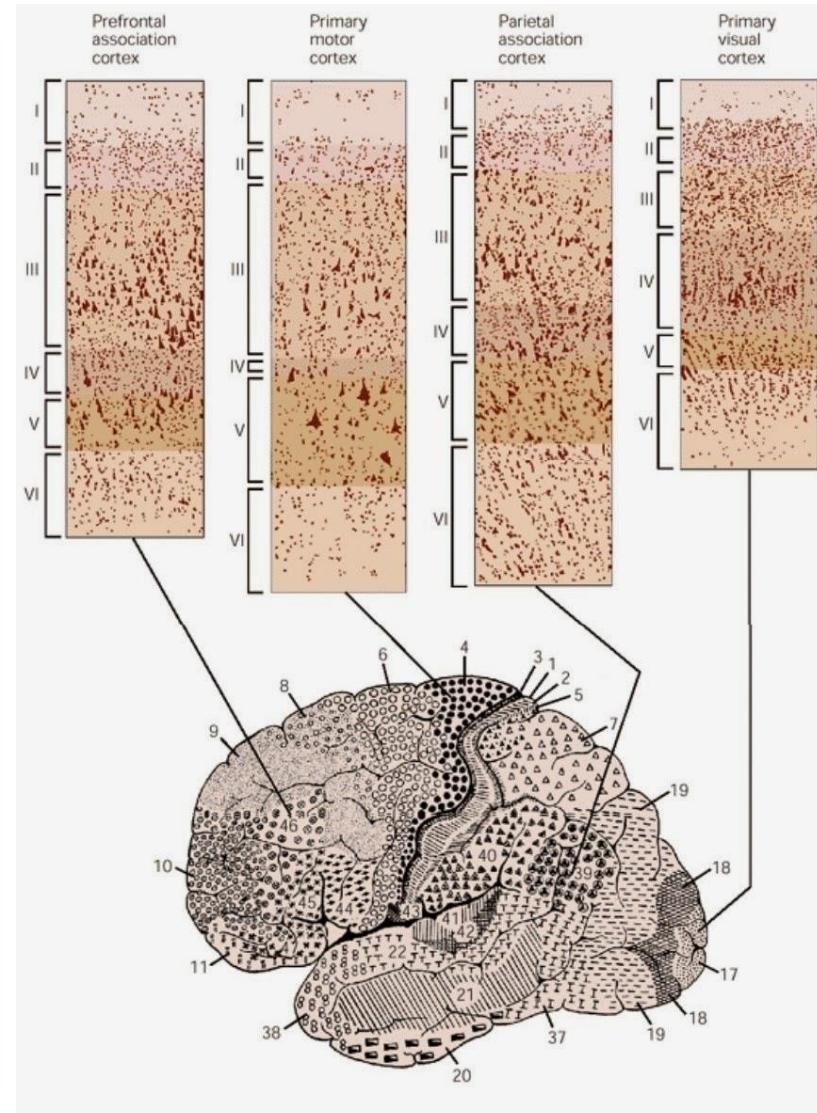
인지, 추론, 운동, 학습 등을 담당하는 대뇌 피질은 구분되는 여러 개의 신경세포 층들로 이루어져 있다.

대뇌피질의 각 영역과, 서로 다른 층은 각기 다른 뇌 기능을 담당하고 있고, 이들은 서로 연결되어 있다.

시각 정보는 Primary Visual Cortex (V1)에서 처리한다.



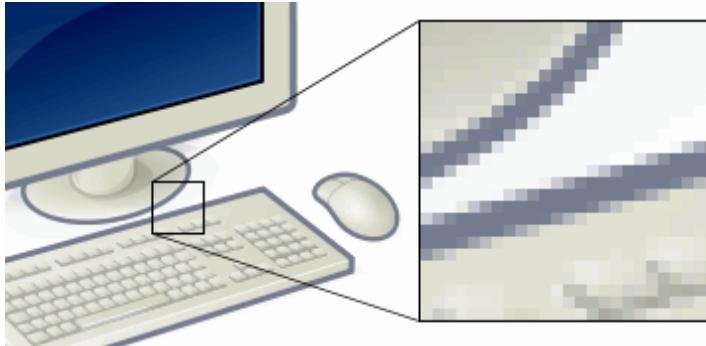
[https://commons.wikimedia.org/wiki/File:Brodmann\\_Cytoarchitectonics\\_17.png](https://commons.wikimedia.org/wiki/File:Brodmann_Cytoarchitectonics_17.png)



<http://theanalogicalinstinct.blogspot.kr/>

# Computer Vision 컴퓨터 비전

픽셀(pixel)로 이루어진 이미지(동영상)로부터 물체의 종류, 위치, 정보, 관계 등 추상적인 정보를 추출하는 작업



<https://en.wikipedia.org/wiki/Pixel>



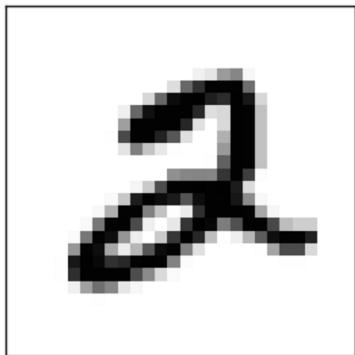
0	2	15	0	0	11	10	0	0	0	0	9	9	0	0	0
0	0	0	4	60	157	236	255	255	177	95	61	32	0	0	29
0	10	16	119	238	255	244	245	243	250	249	255	222	103	10	0
0	14	170	255	255	254	254	255	253	245	255	249	253	251	124	1
2	98	255	228	255	251	254	211	141	116	122	215	251	238	255	49
13	217	243	255	155	33	226	52	2	0	10	13	232	255	255	36
16	229	252	254	49	12	0	7	7	0	70	237	252	235	62	6
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0
0	87	252	250	248	215	60	0	1	121	252	255	248	144	6	0
0	13	119	255	255	245	255	182	181	248	252	242	208	36	0	19
1	0	5	117	251	255	241	255	247	255	241	162	17	0	7	0
0	0	0	4	58	251	255	246	254	253	255	120	11	0	1	0
0	0	4	9	255	255	255	248	252	255	244	255	182	10	0	4
0	22	206	252	246	251	241	100	24	115	255	245	255	194	9	0
0	111	255	242	255	158	24	0	0	6	39	255	232	230	56	0
0	218	251	250	137	7	11	0	0	0	62	255	250	125	3	0
0	173	255	255	101	9	20	0	13	3	13	182	251	245	61	0
0	107	251	241	255	230	98	55	19	115	217	248	253	255	52	4
0	18	146	250	255	247	255	255	249	255	240	255	129	0	5	0
0	0	23	115	215	255	250	248	255	255	248	248	118	14	12	0
0	0	6	1	0	52	153	233	255	252	147	37	0	0	4	1
0	0	5	5	0	0	0	0	14	1	0	6	6	0	0	0

<https://mozanunal.com/2019/11/img2sh/>

# Visual Recognition. 시각 인지(인식)

이미지(숫자들의 배열)이 나타내는 대상의 종류를 파악하는 지능

From an array of numbers to a category



2



# Visual Recognition. 시각 인지(인식)

이미지(숫자들의 배열)이 나타내는 대상의 종류를 파악하는 지능

From an array of numbers to a category

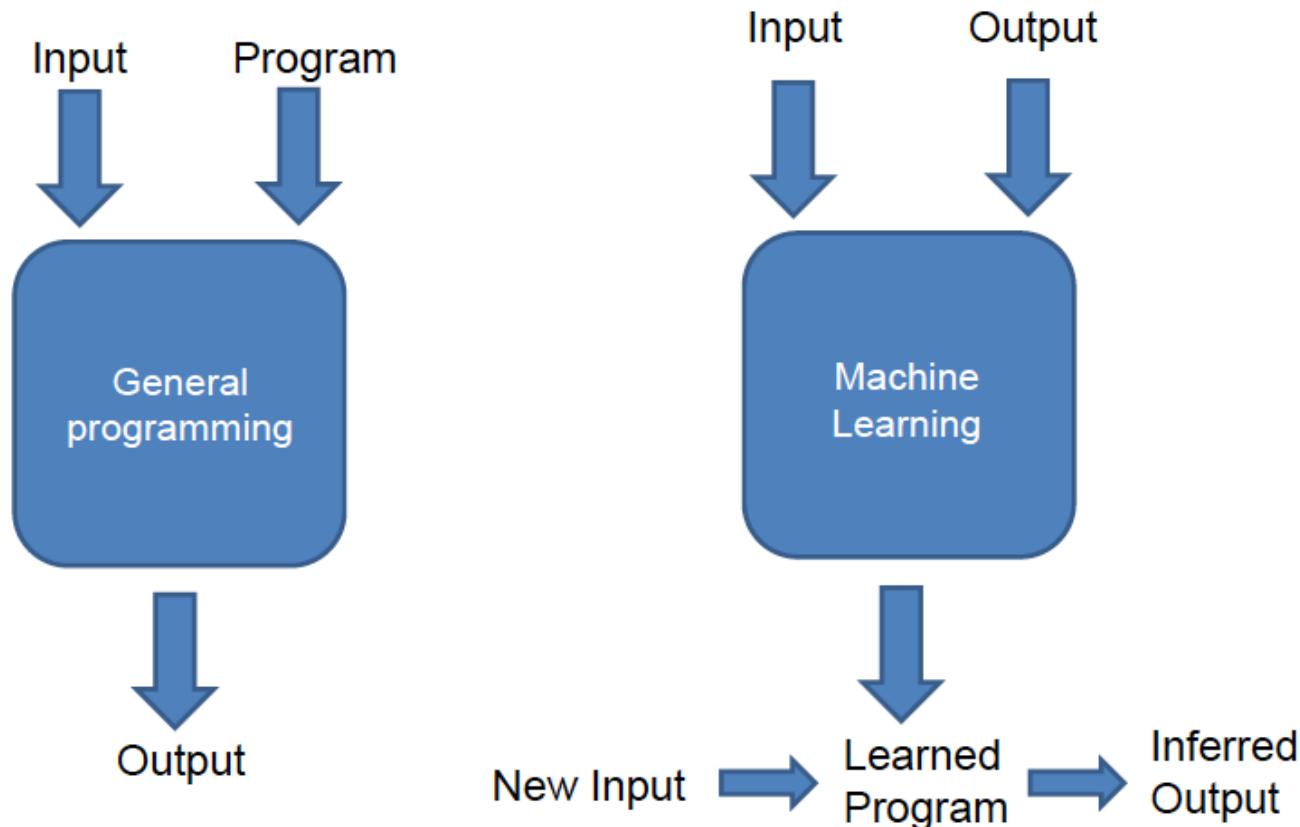


<https://www.pngkey.com/maxpic/u2q8t4t4u2e6w7e6/>

# MNIST: Dataset for Handwritten Digit Recognition

# Machine Learning 기계학습

## General Programming vs Machine Learning

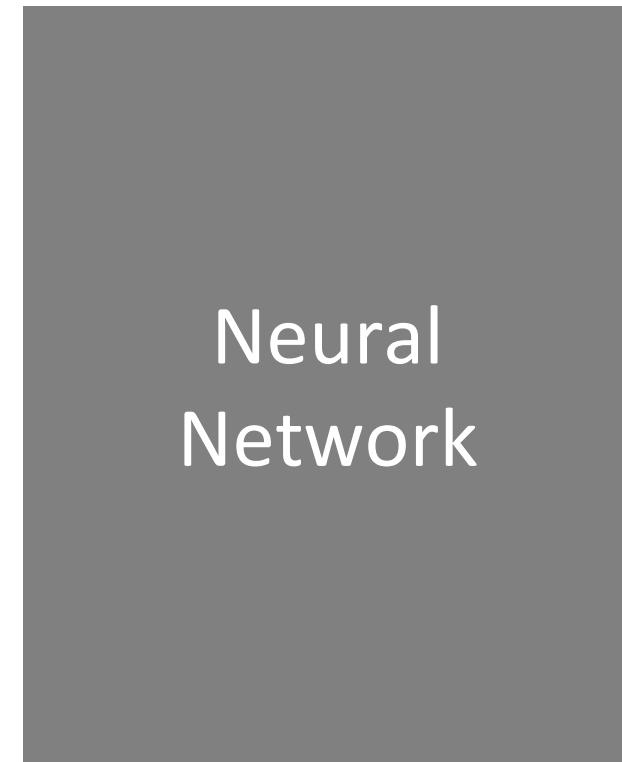
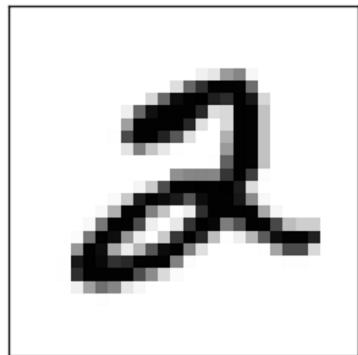


Ibishal, Eindhoven Developers Meetup, 2015

# Visual Recognition. 시각 인지(인식)

이미지(숫자들의 배열)이 나타내는 대상의 종류를 파악하는 지능

From an array of numbers to a category

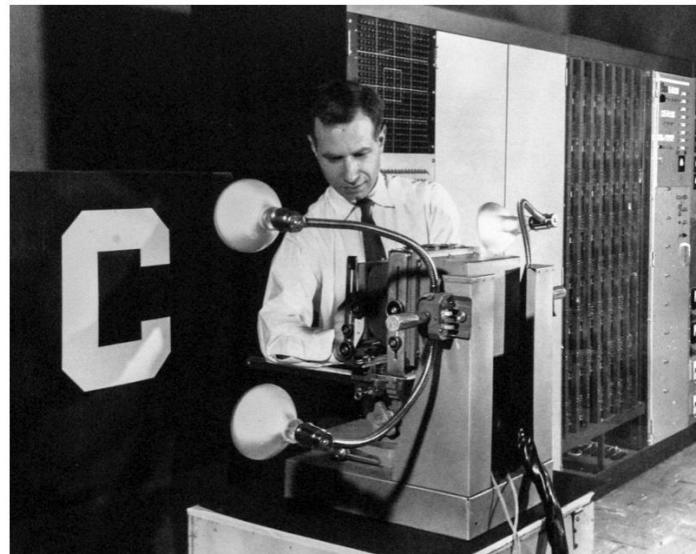
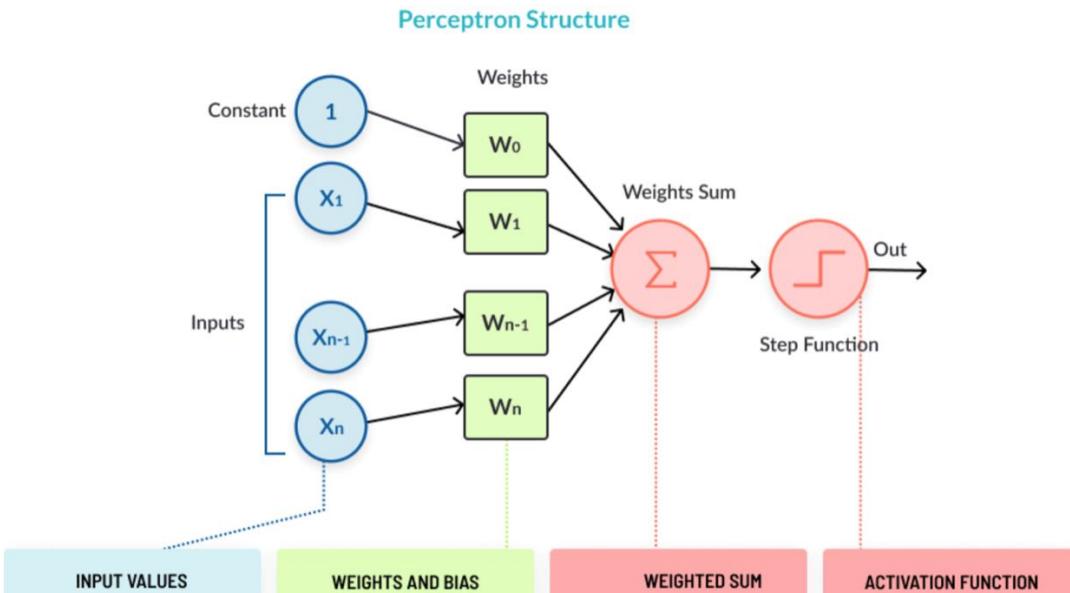


2



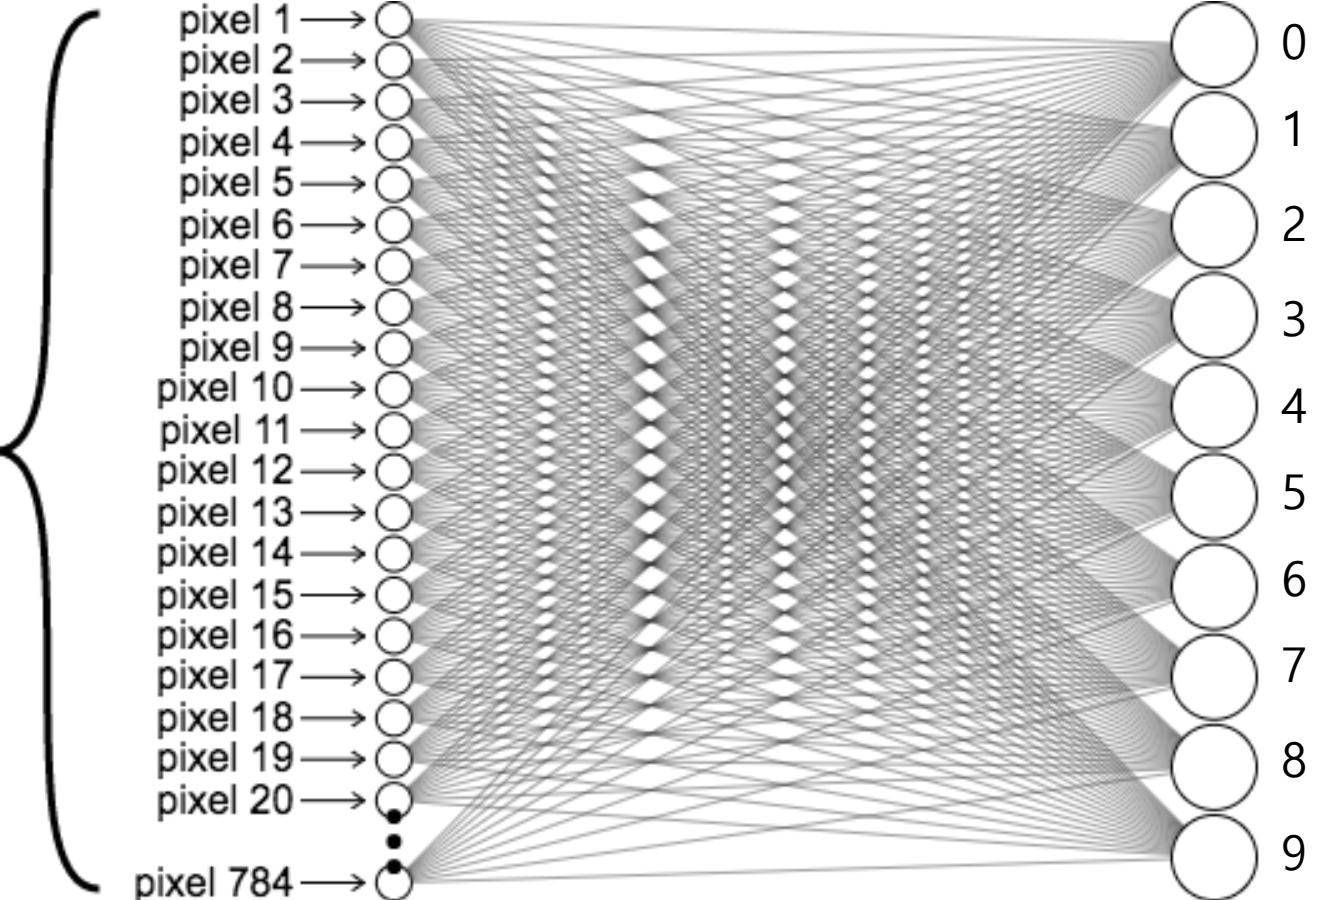
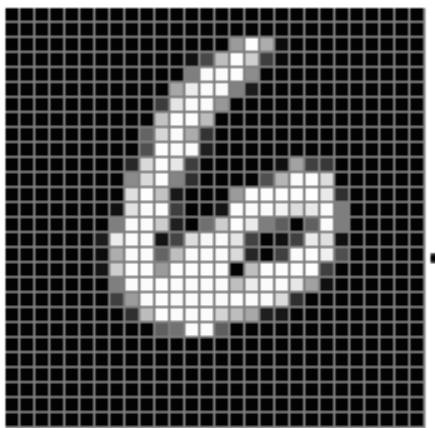
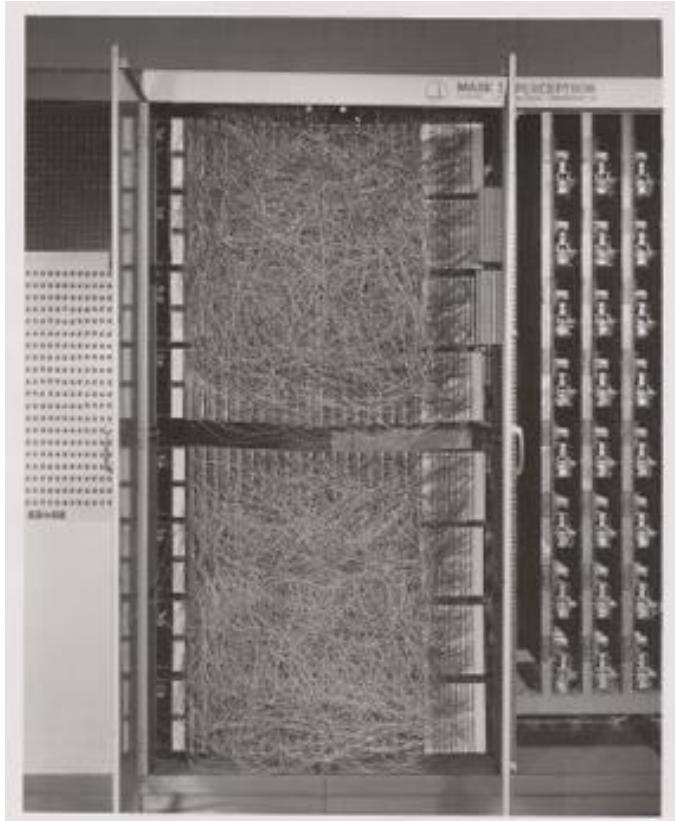
# Mark I Perceptron

## 1st Generation: Binary Neurons



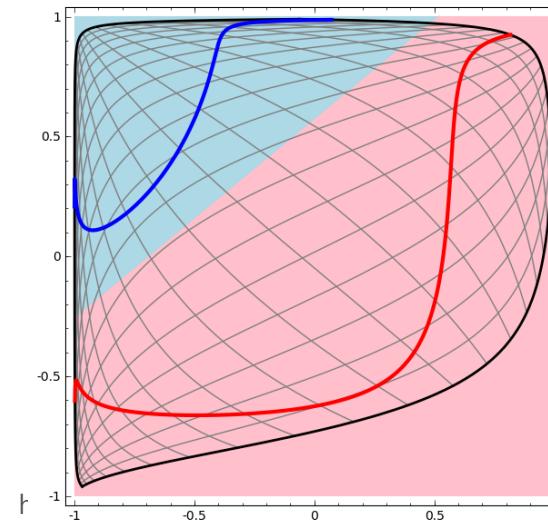
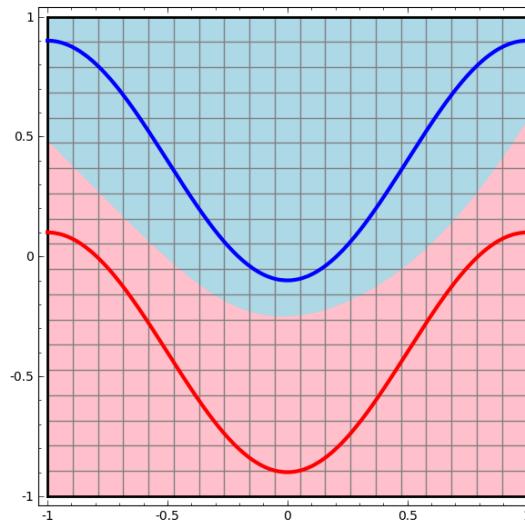
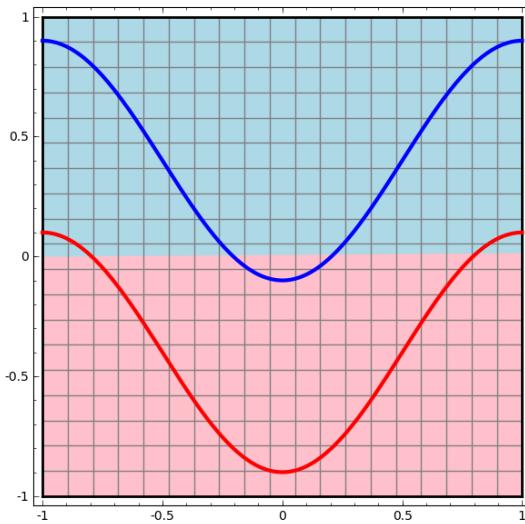
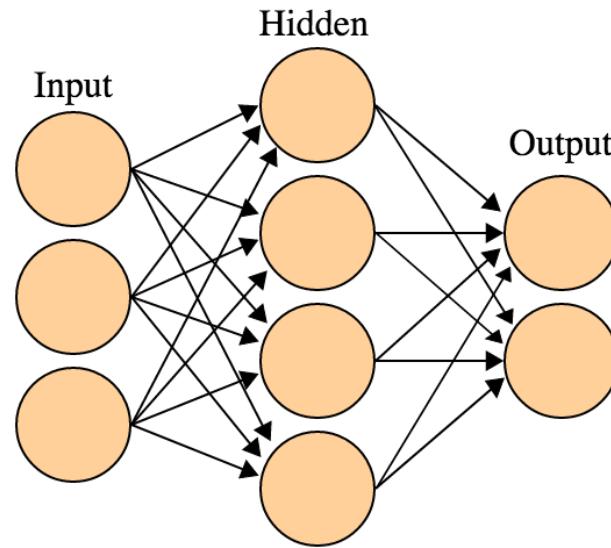
<https://news.cornell.edu/stories/2019/09/professors-perceptron-paved-way-ai-60-years-too-soon>

# Mark I Perceptron

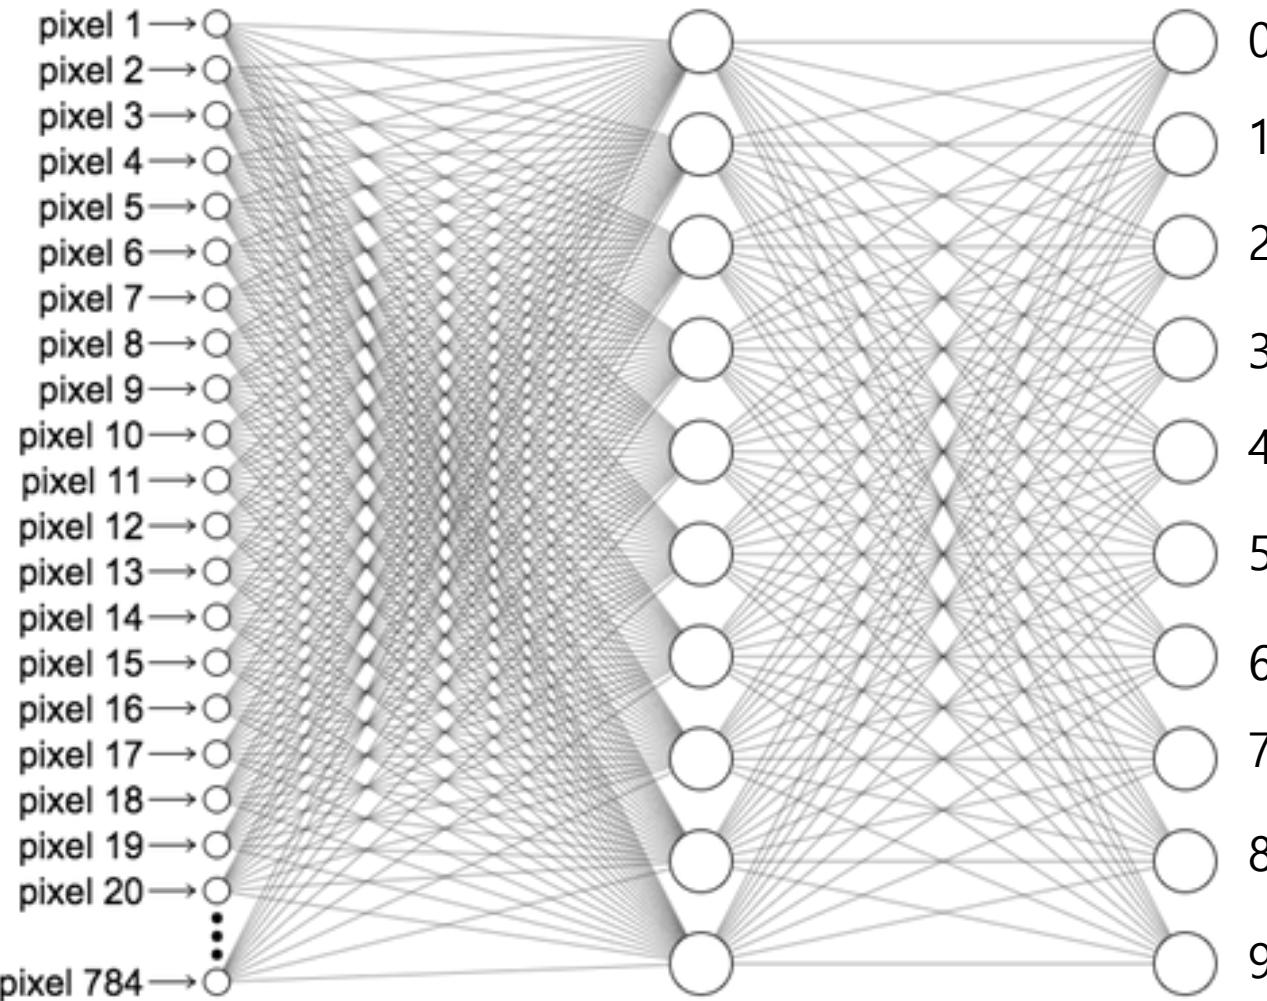
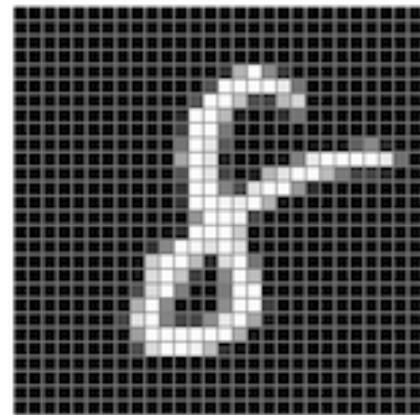


<https://digital.library.cornell.edu/catalog/ss:550351>

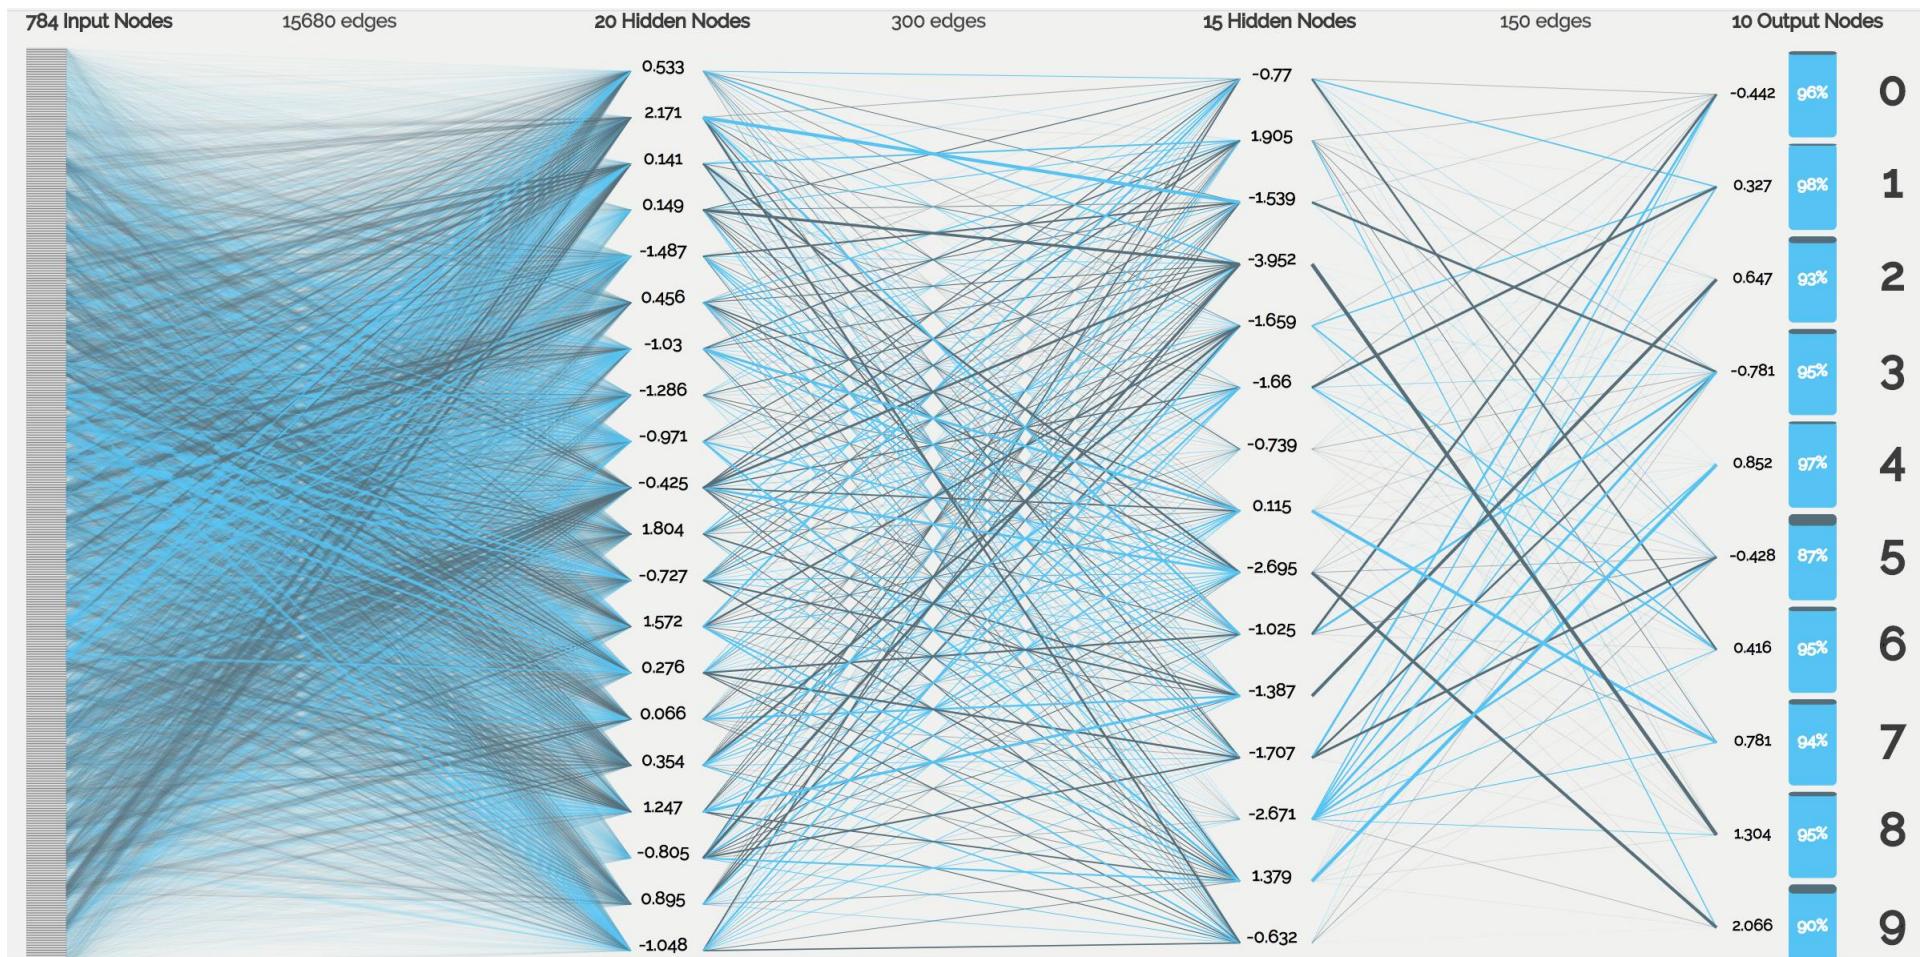
# Multilayer Perceptron



# Multilayer Perceptron or Fully-connected Neural Network

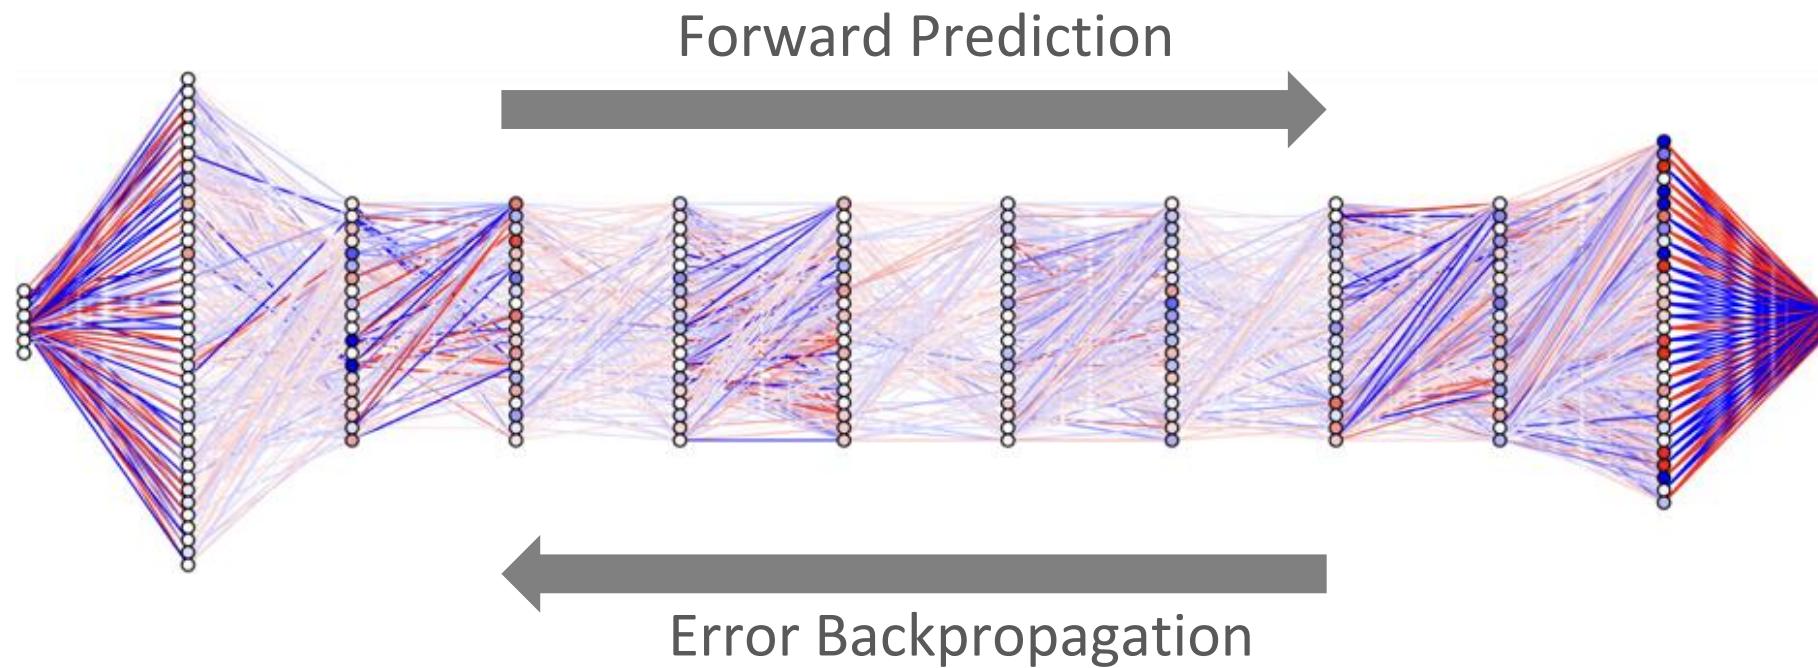


# Multilayer Perceptron or Fully-connected Neural Network



<http://nn-mnist.sennabaum.com/>

# Deep Neural Networks with Backpropagation



<https://www.mghpcc.org/neural-networks-earthquakes/>

# Moravec's Paradox

## Moravec Paradox & The nature of intelligence



Kasparov vs. machine (2003)

Fake victory of AI (brut force computation)

“The main lesson of thirty-five years of AI research is that. **the hard problems are easy and the easy problems are hard**”

Linguist and cognitive scientist *Steven Pinker*

“It is easy to make computers exhibit adult level performance on intelligence tests or playing checkers, and **difficult or impossible to give them the skills of a one-year-old** when it comes to perception and mobility” - 1988

Moravec



# Challenges of Visual Recognition

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter



Intra-class variation

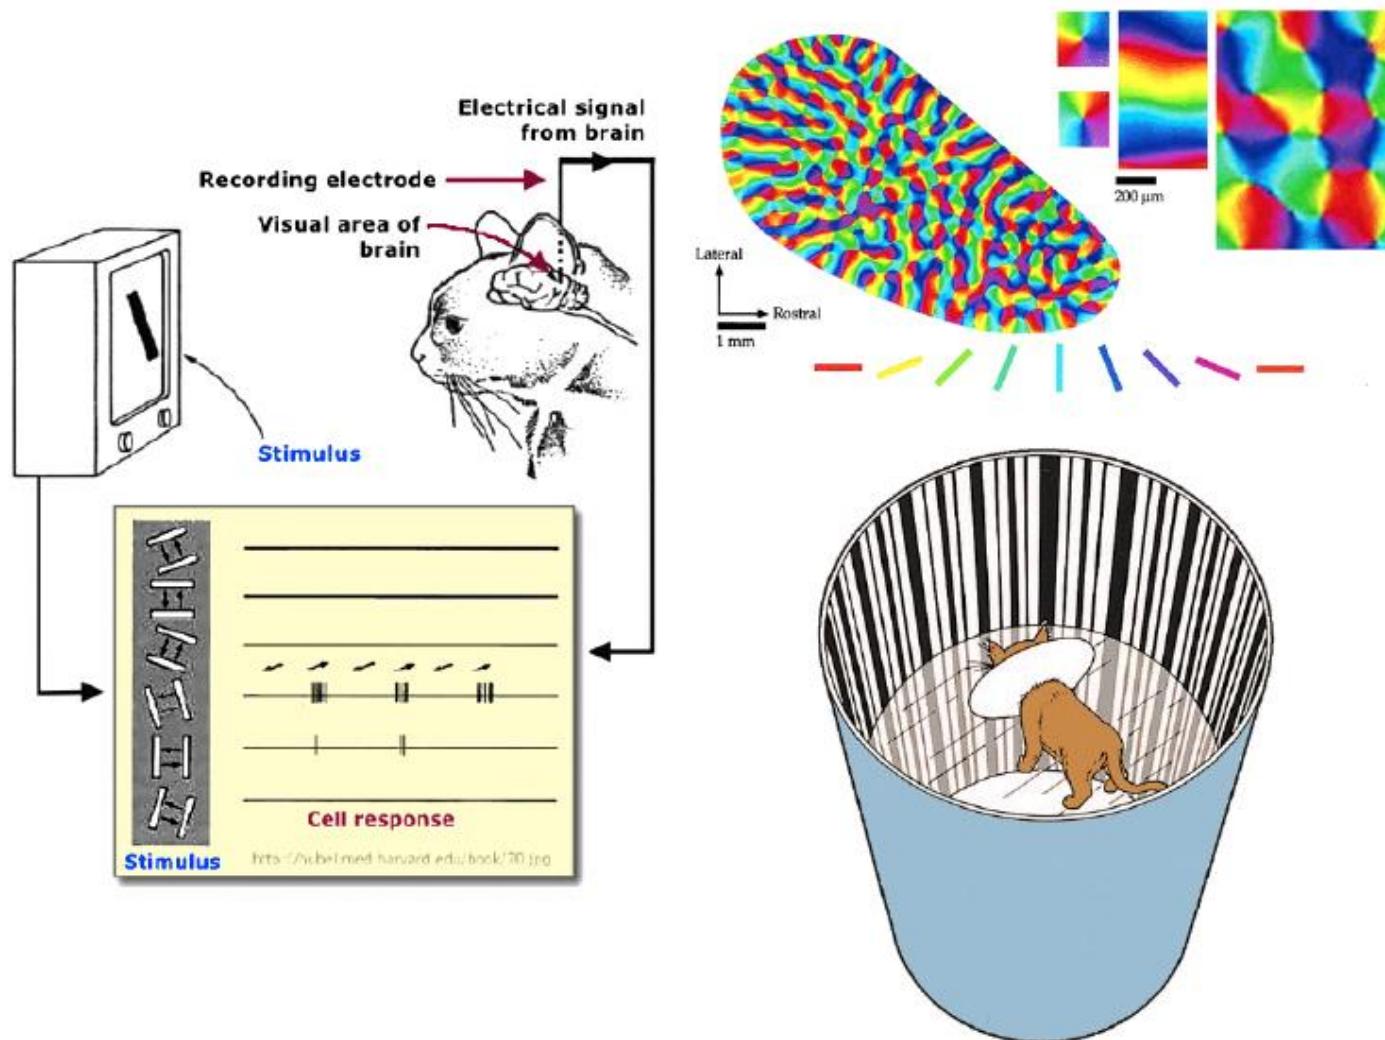


## Hint from Neuroscience: The Hubel and Wiesel Experiment



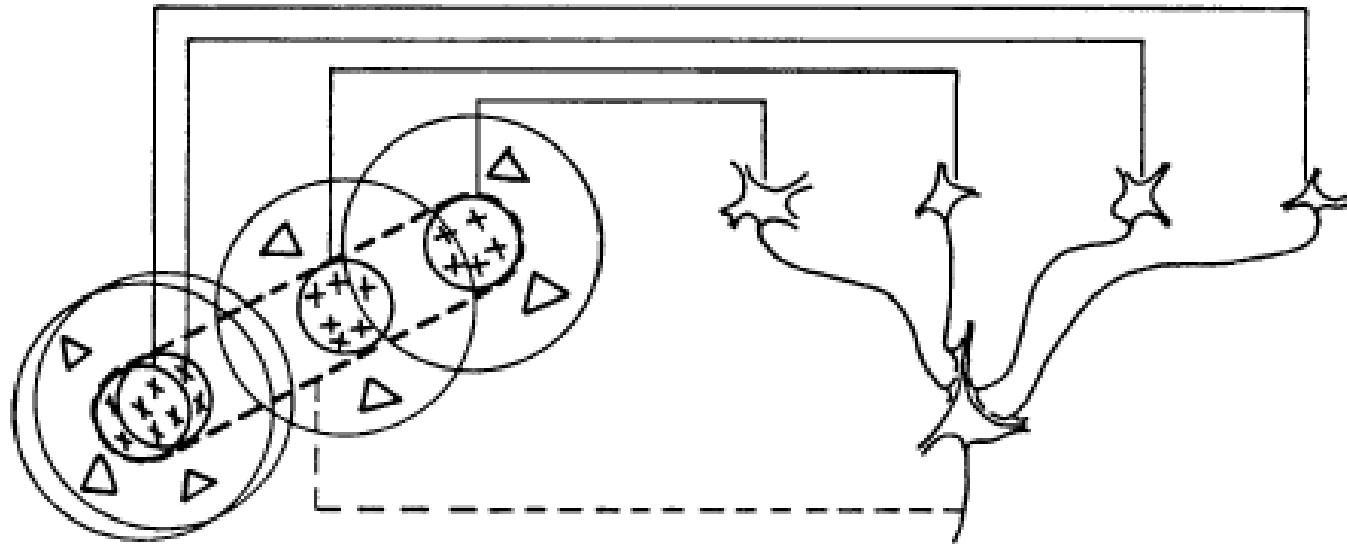
<https://www.youtube.com/watch?v=IOHayh06LJ4>

# Neurons responding to a certain visual pattern



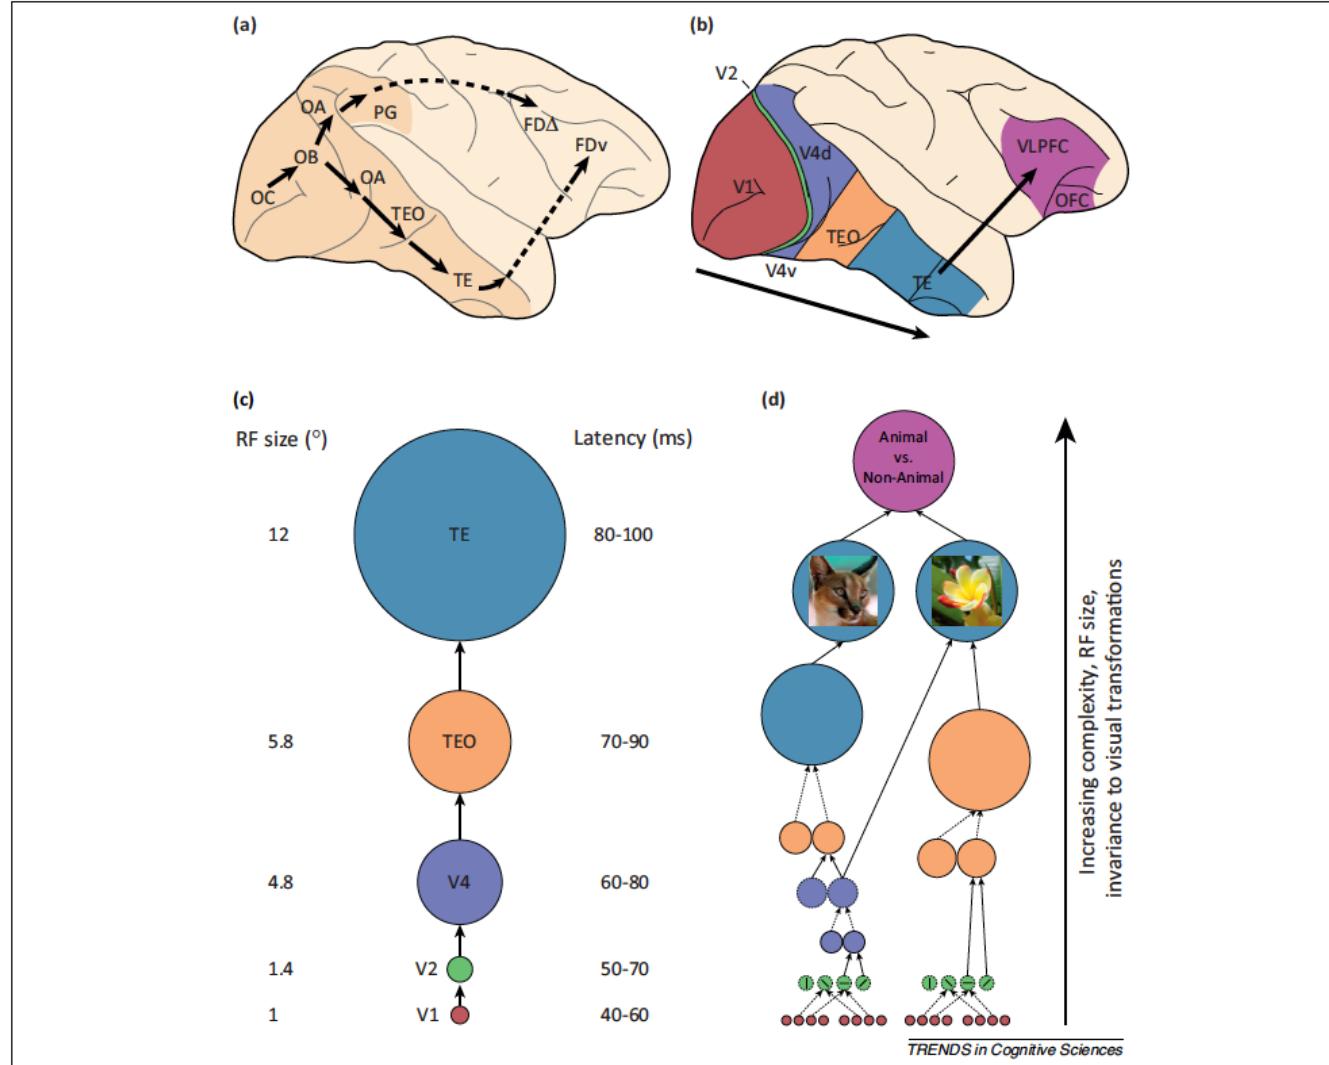
Oxford Machine Learning Course

# Receptive Fields



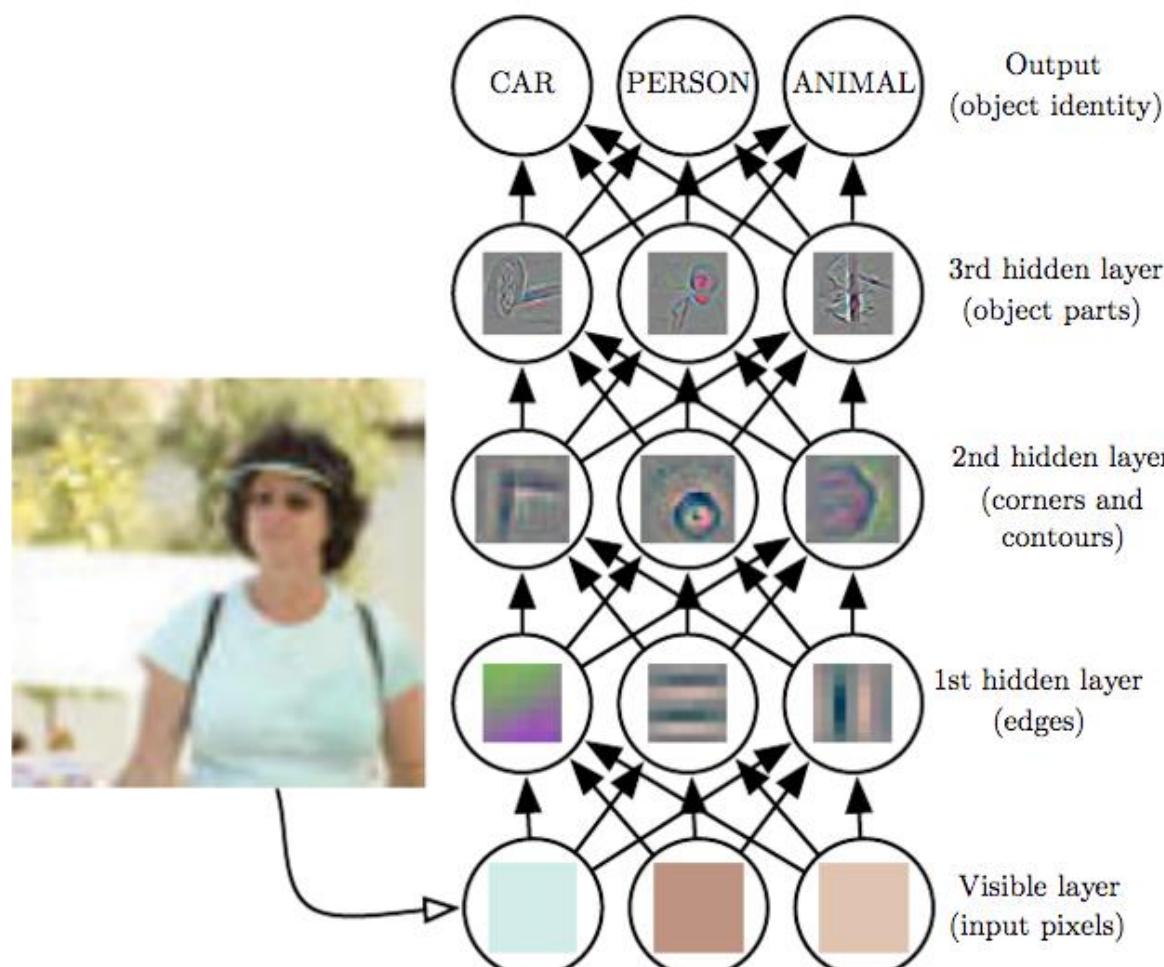
Hubel and Wiesel, *RECEPTIVE FIELDS, BINOCULAR INTERACTION AND FUNCTIONAL ARCHITECTURE IN THE CAT'S VISUAL CORTEX*, Journal of Physiology (1962)

# Hierarchical Organization of Visual Cortex



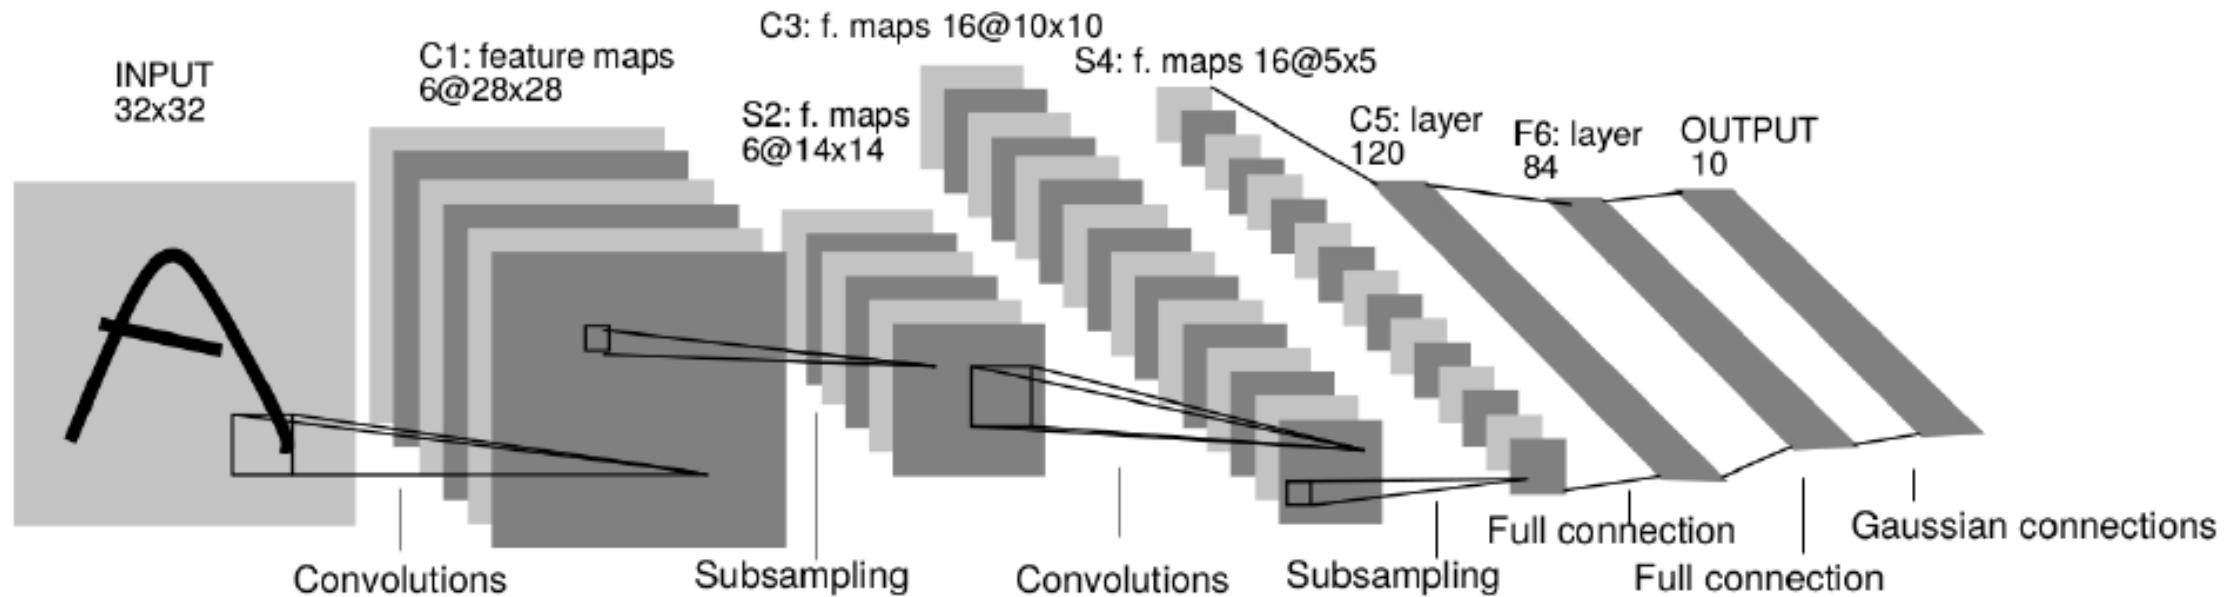
Kravitz et al., "The ventral visual pathway: an expanded neural framework for the processing of object quality", *Trends in Cognitive Sciences*, 2013

# Representation Learning for Visual Recognition



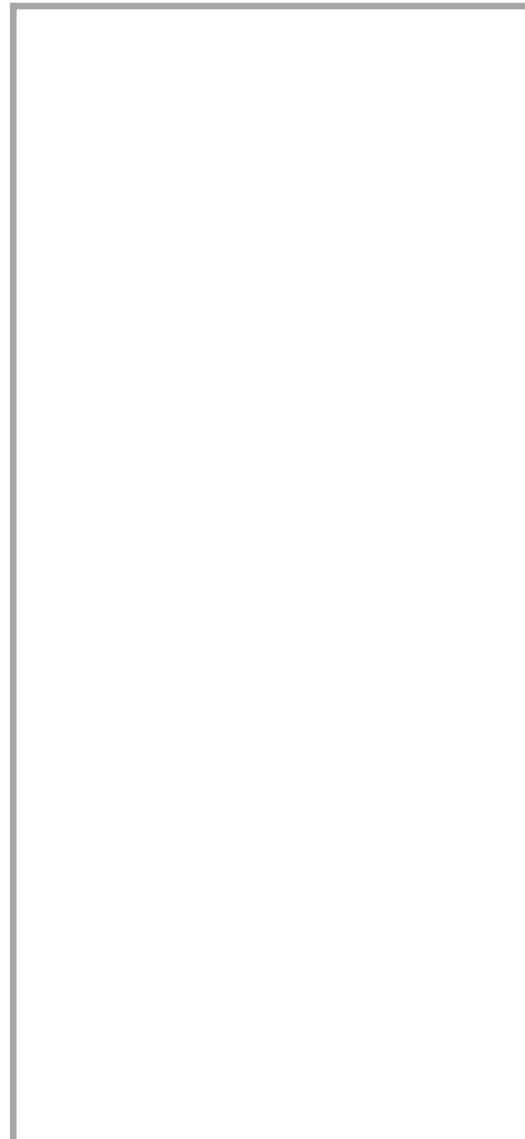
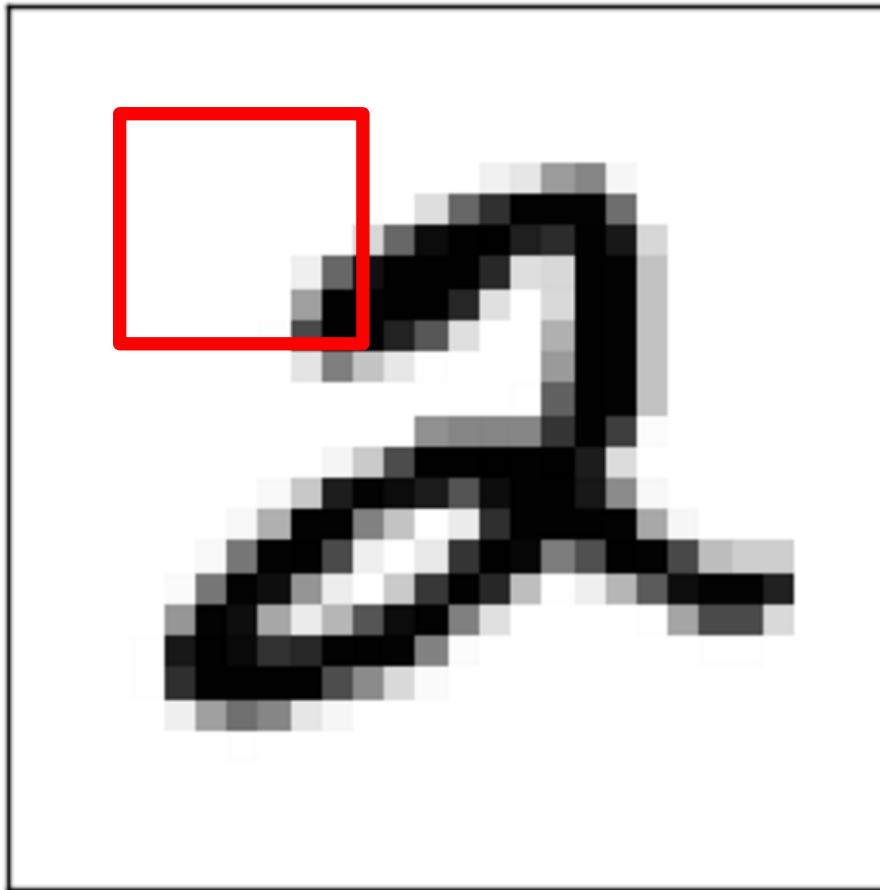
Goodfellow et al., 2015

# Convolutional Neural Network: LeNet-5

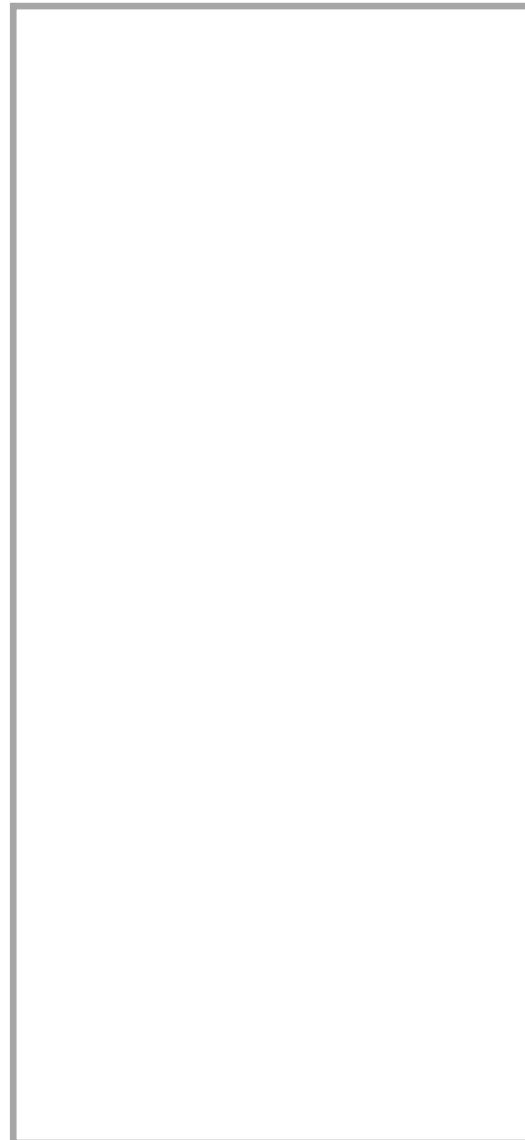
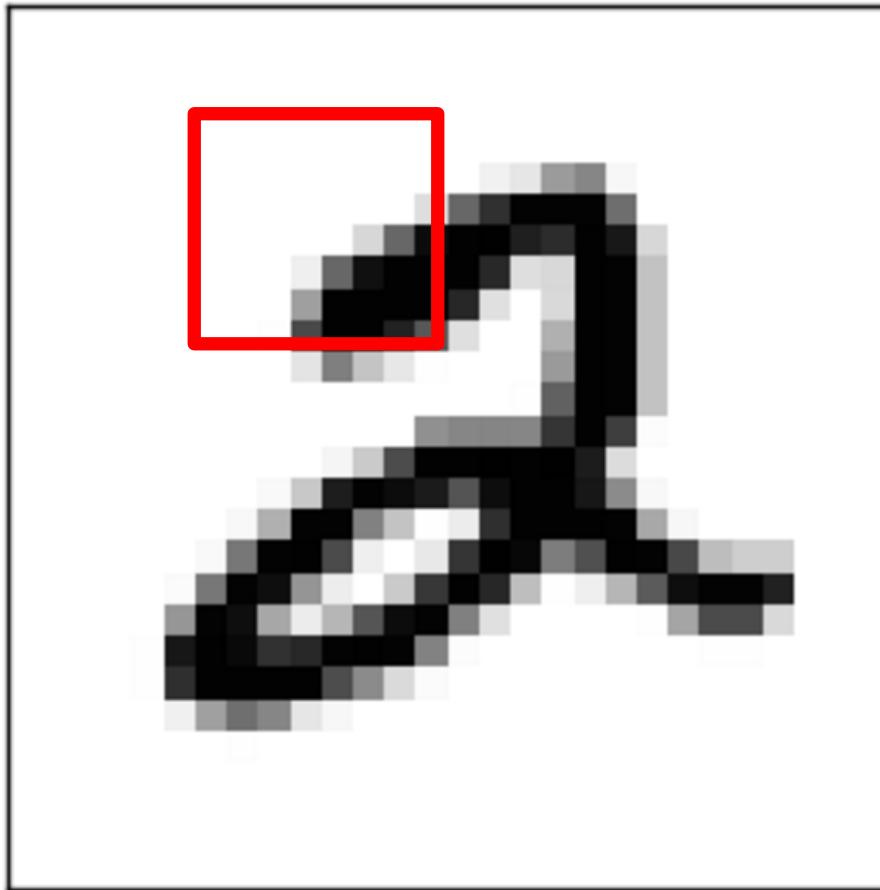


LeCun et al., Gradient-based learning applied to document recognition. Proceedings of the IEEE 86, 11 (1998)

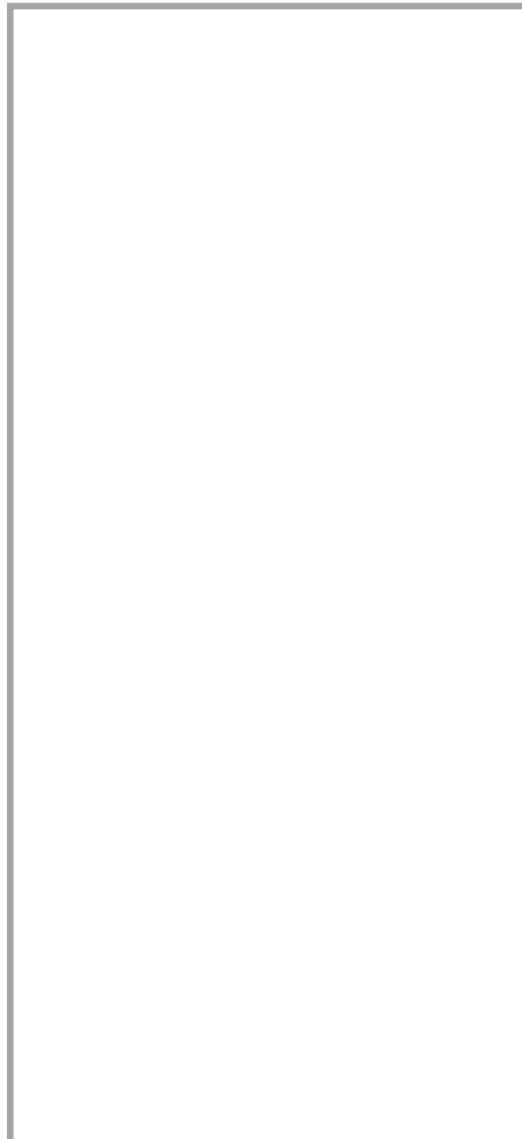
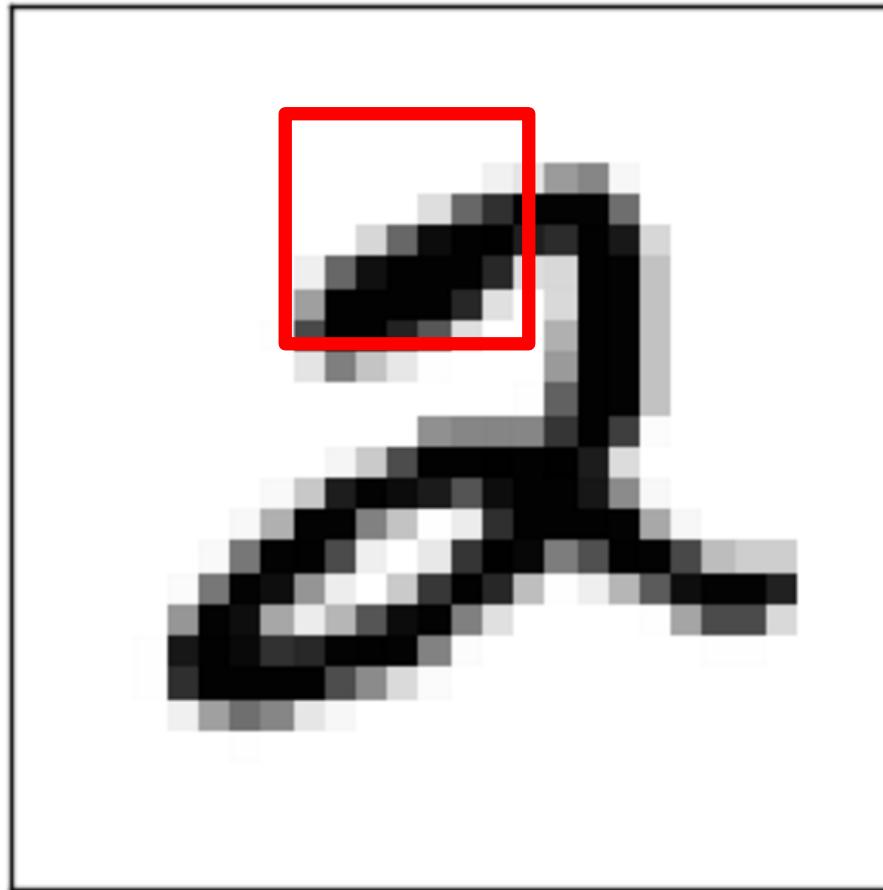
# Convolution Operation (2D)



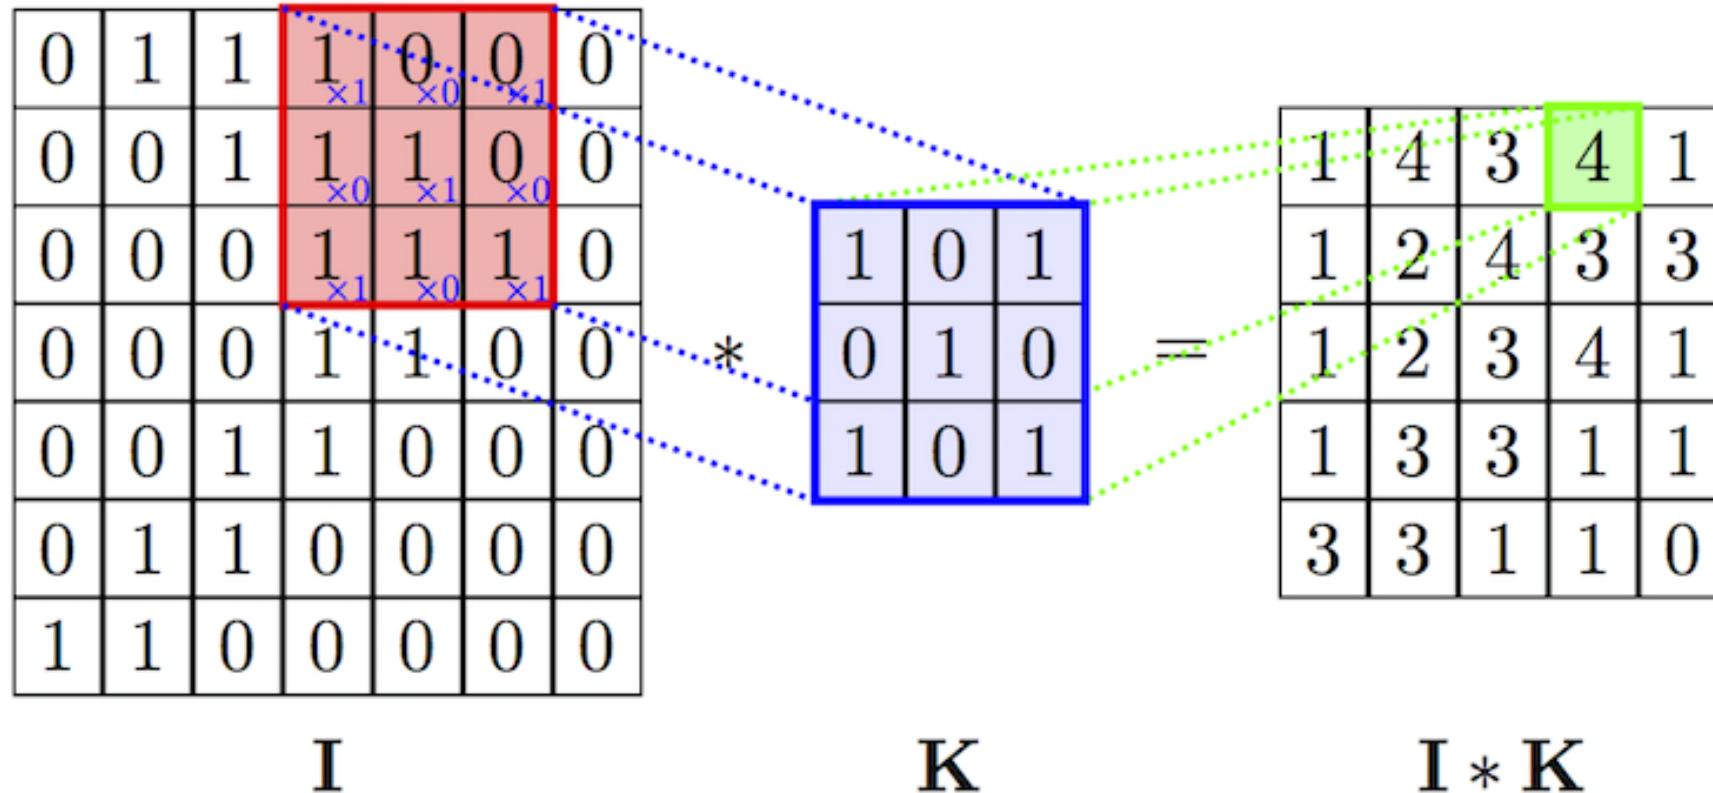
# Convolution Operation (2D)



# Convolution Operation (2D)

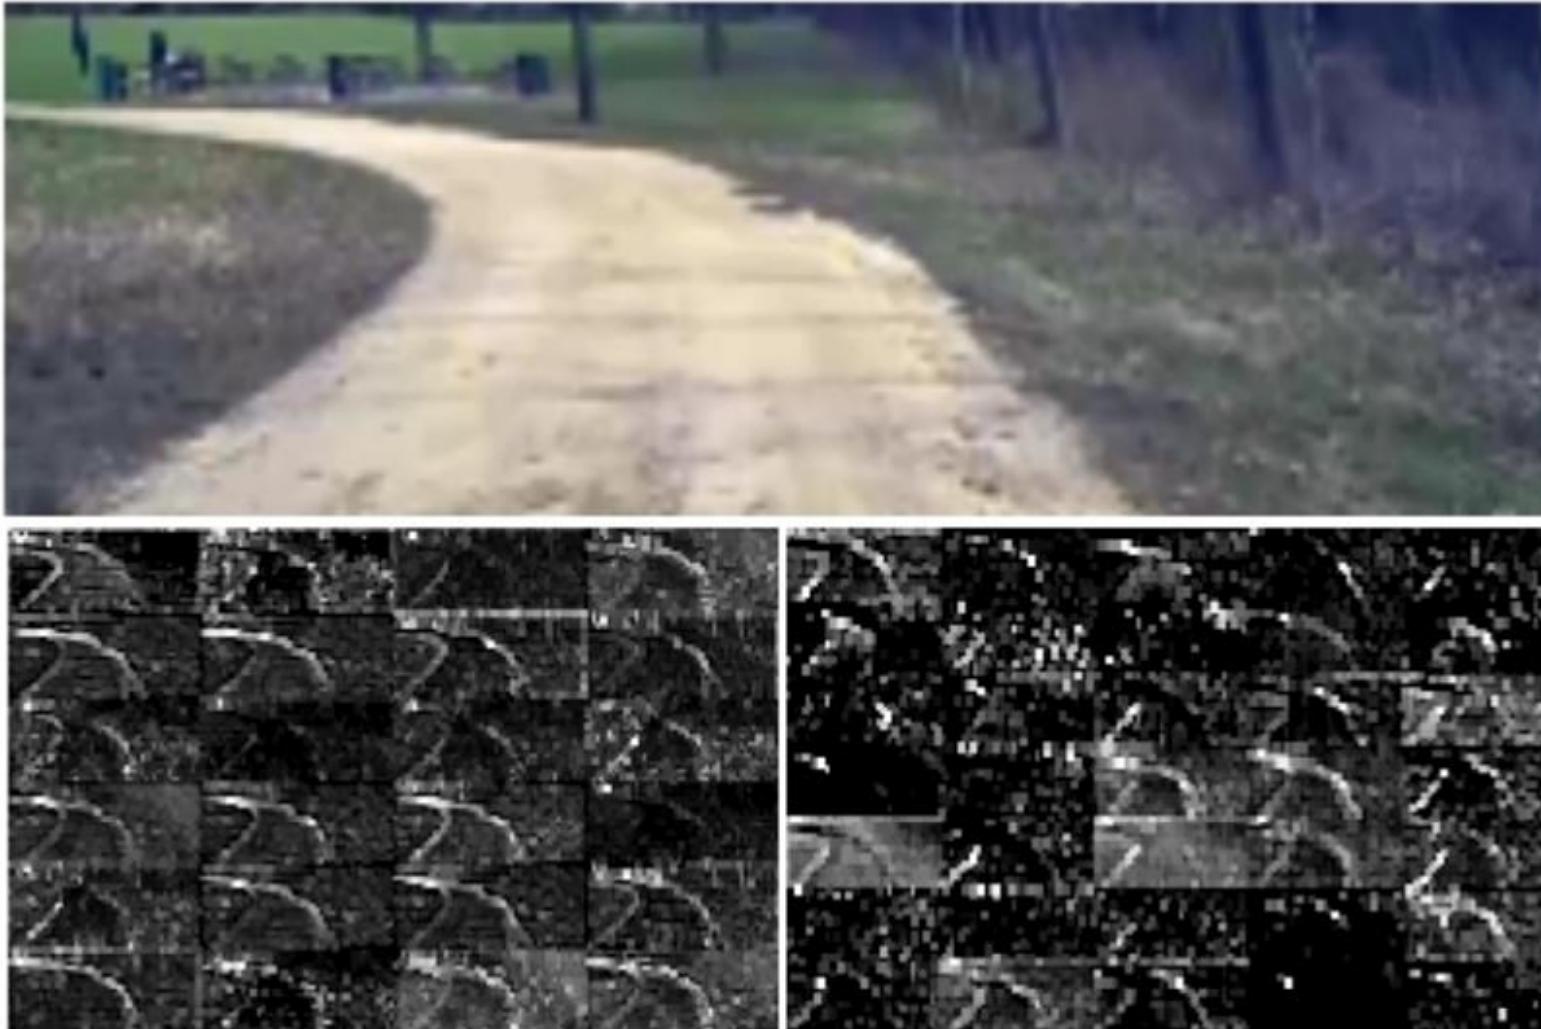


# Convolution Operation (2D)



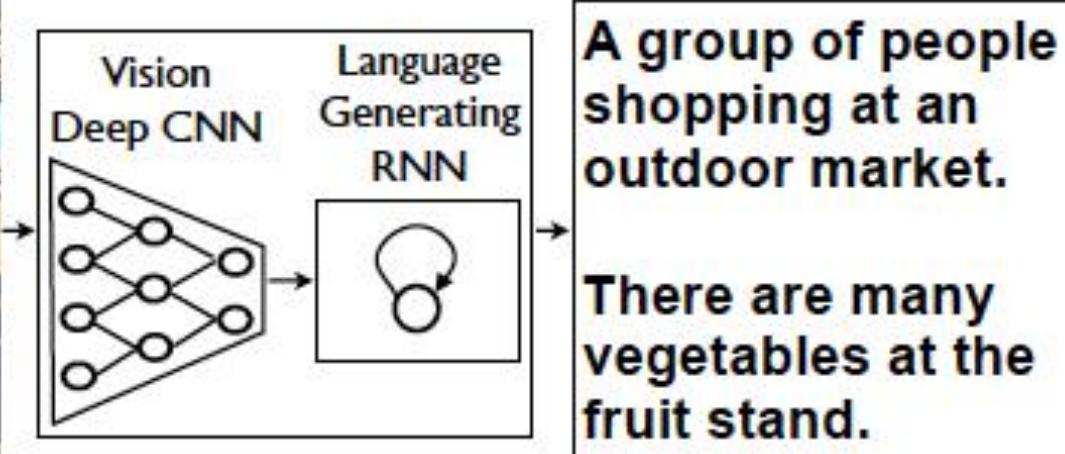
<https://cambridgespark.com/>

# Self-driving Cars



Bojarski et al., "End-to-End Learning for Self-Driving Cars", *arXiv preprint*, 2016

# Image Captioning



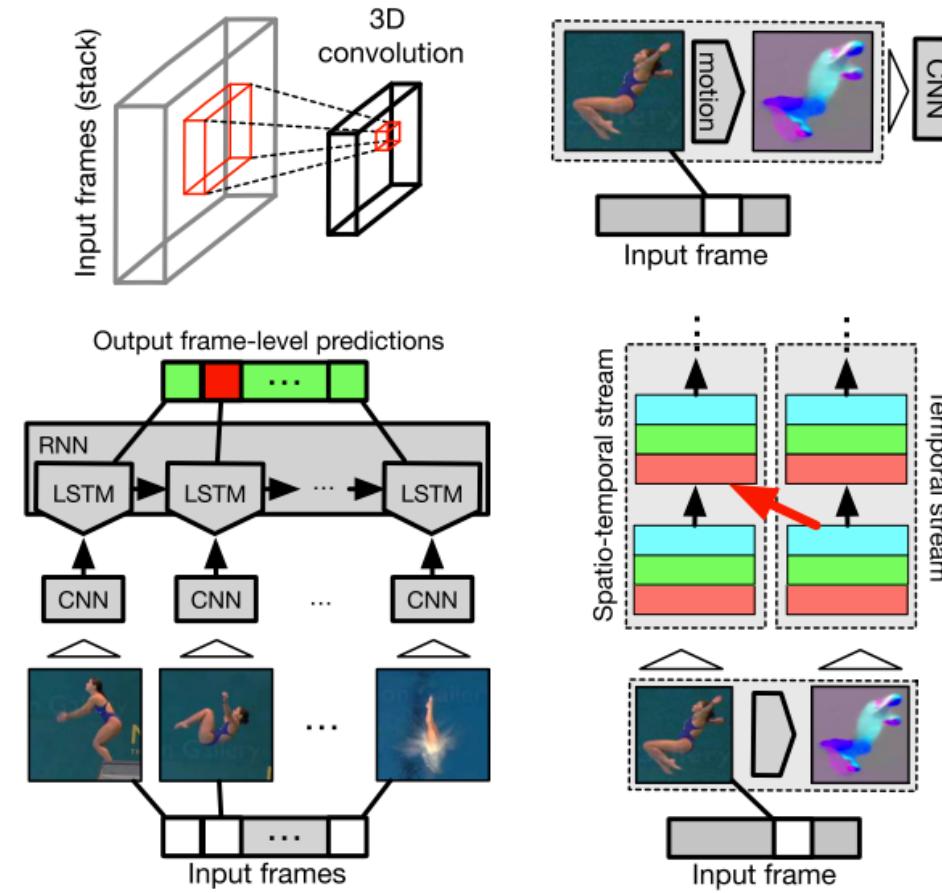
Vinyals et al., *Show and Tell: A Neural Image Caption Generator*, arXiv (2015)

# Video Understanding

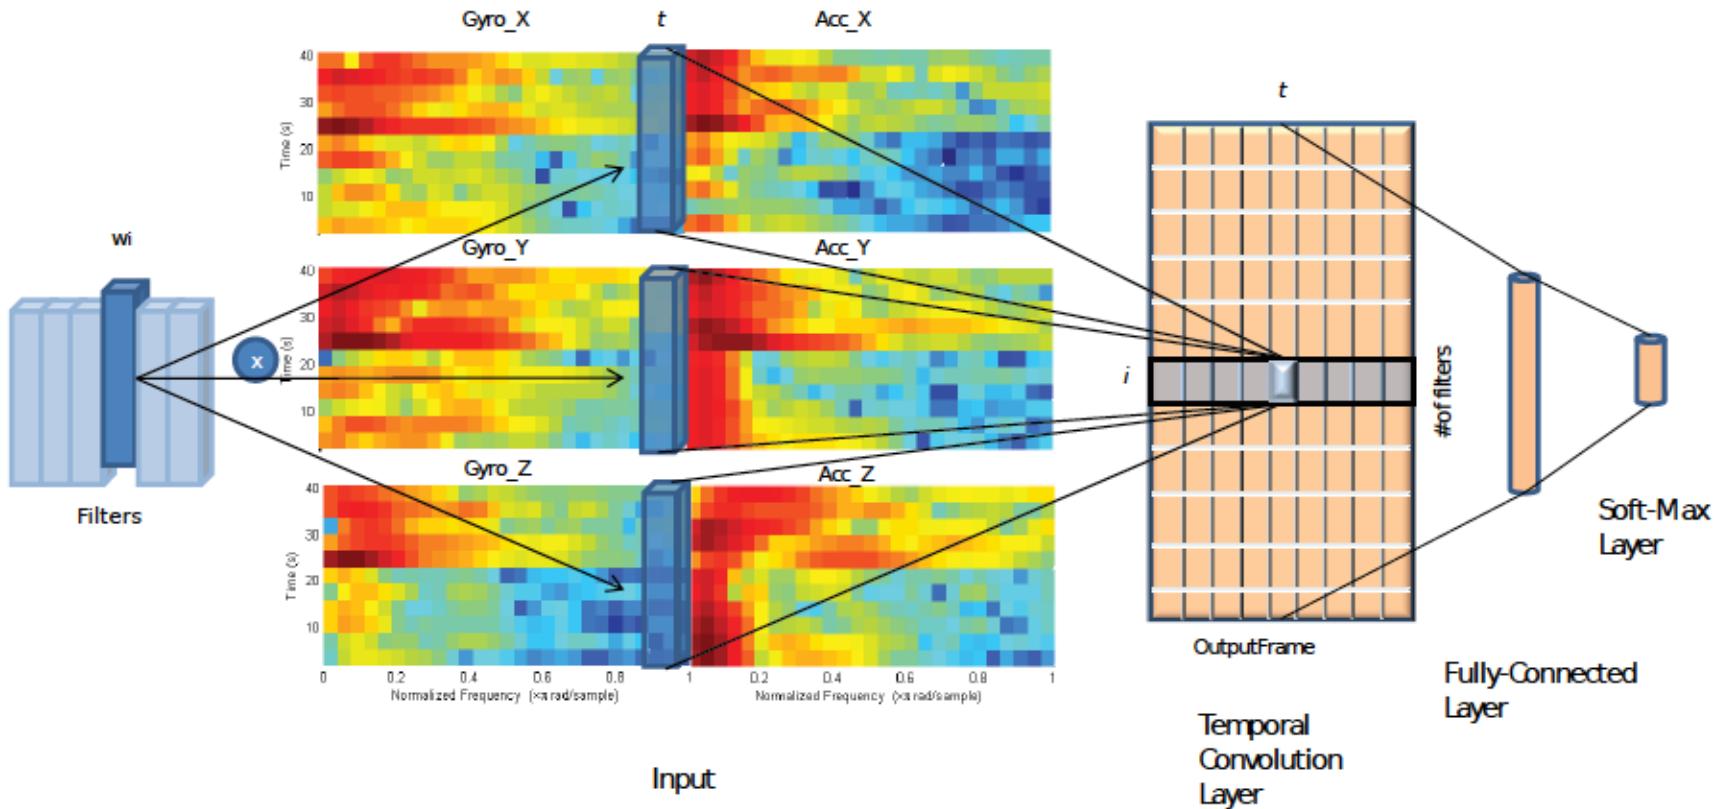
2017 IEEE 12th International Conference on Automatic Face & Gesture Recognition

## A survey on deep learning based approaches for action and gesture recognition in image sequences

Maryam Asadi-Aghbolaghi<sup>1,2,3</sup>, Albert Clapés<sup>2,3</sup>, Marco Bellantonio<sup>4</sup>, Hugo Jair Escalante<sup>5</sup>, Víctor Ponce-López<sup>2,3,6</sup>, Xavier Baró<sup>6</sup>, Isabelle Guyon<sup>7</sup>, Shohreh Kasaei<sup>1</sup>, Sergio Escalera<sup>2,3</sup>



# Temporal Convolution



Ravi et al., 2016

# Visual Question Answering (VQA)

Who is wearing glasses?

man



woman



Is the umbrella upside down?

yes



no



Where is the child sitting?

fridge



arms



How many children are in the bed?

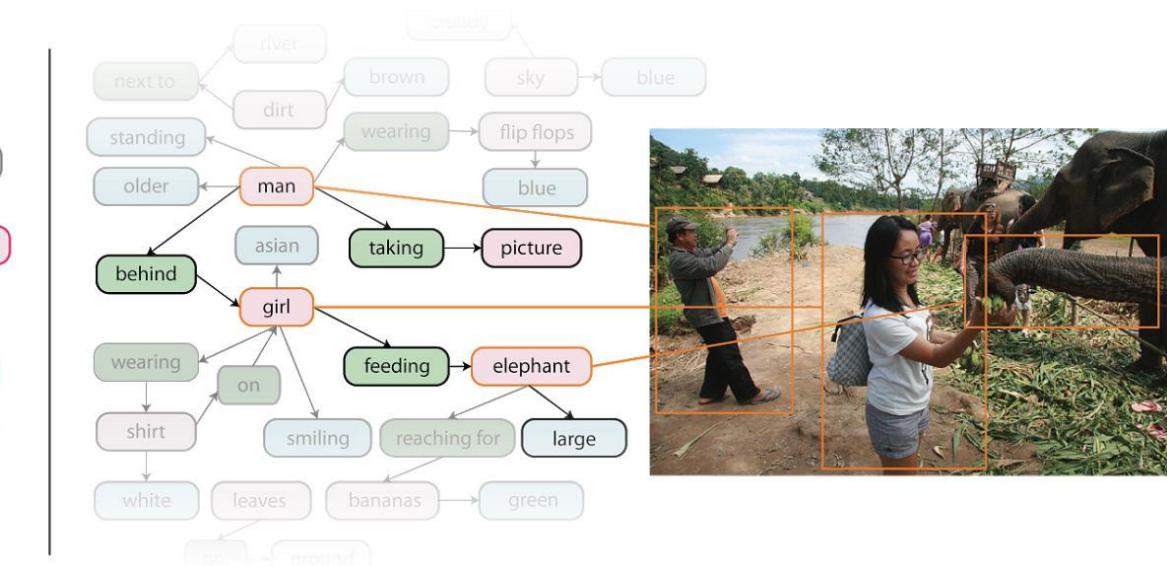
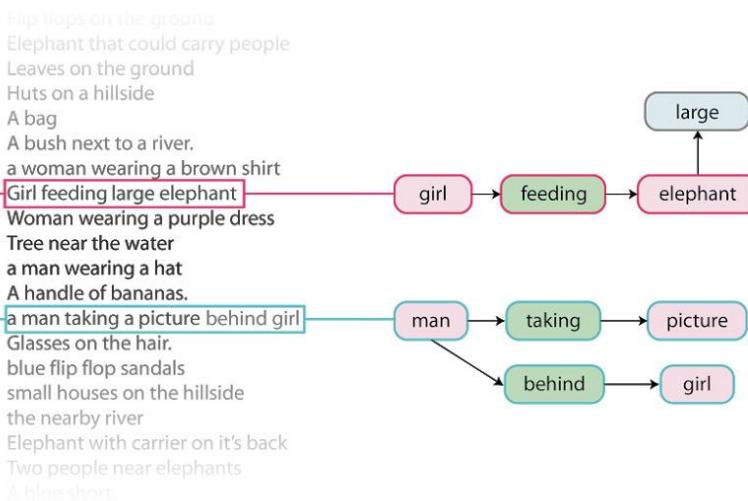
2



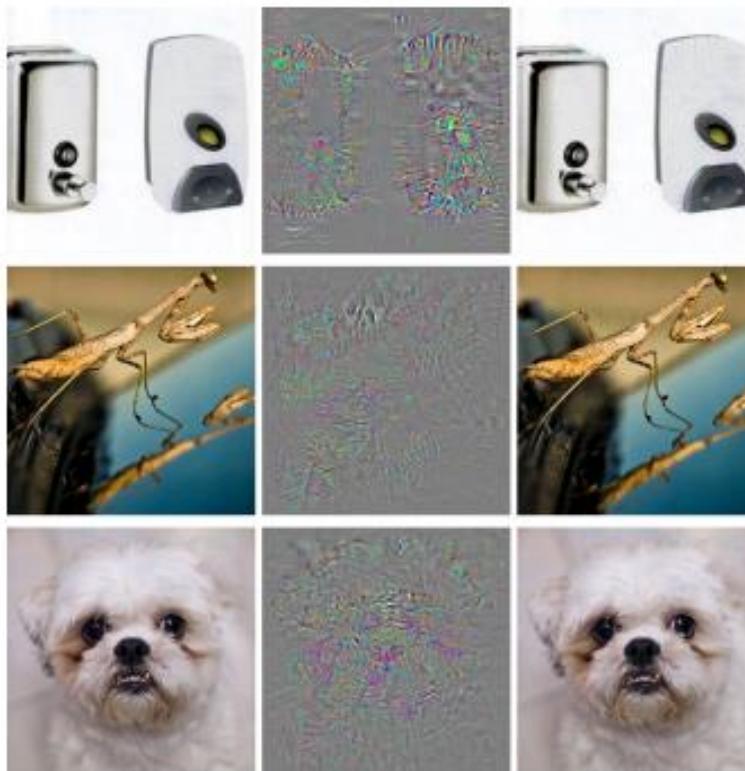
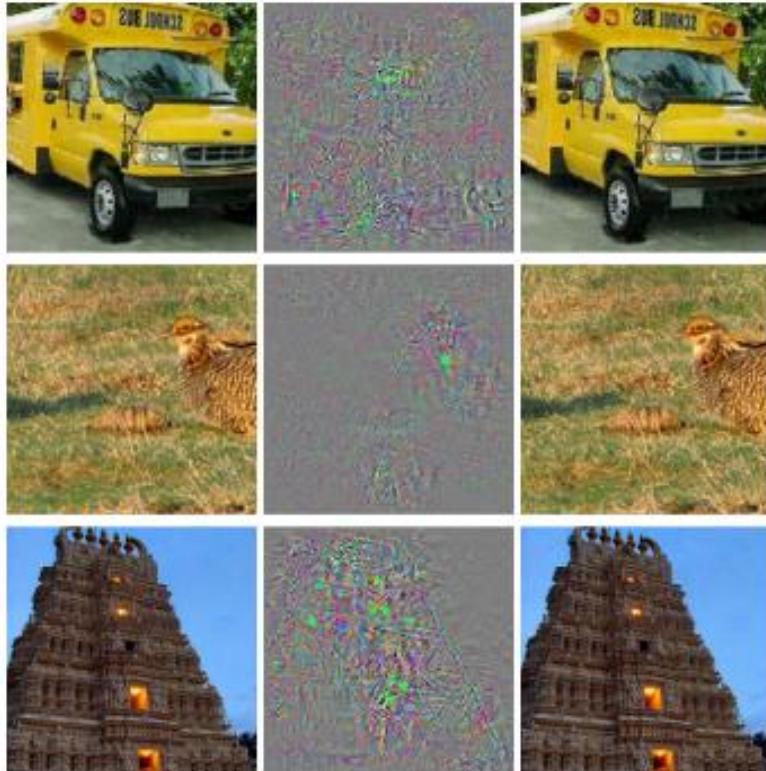
1



# Scene Graph

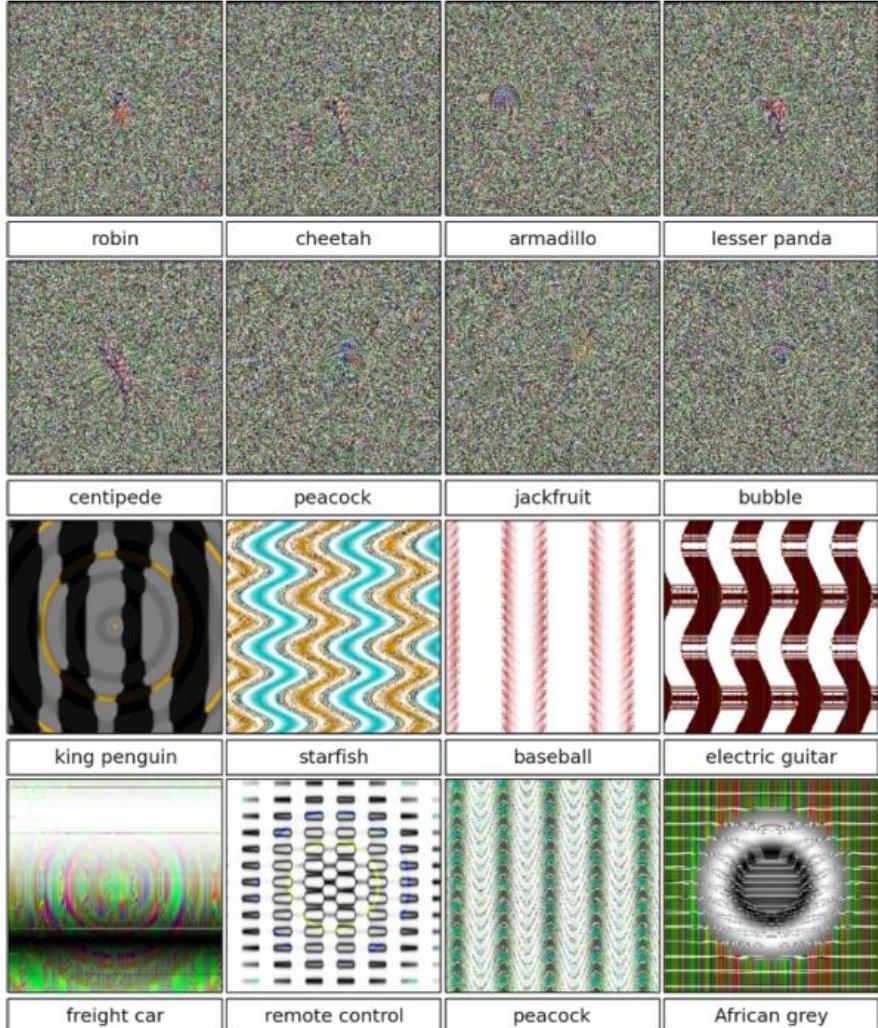


# Fooling Deep Neural Networks



Szegedy et al., "Intriguing Properties of Neural Networks", *arXiv preprint*, 2014

# Fooling Deep Neural Networks



Nguyen et al., "Deep Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images", CVPR, 2015

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Original Series	1	0	2	3	5	1	2	3
Kernel	-1	2	1					
Convolution Series	1							

$$(1, 0, 2) \cdot (-1, 2, 1) = 1*(-1) + 0*2 + 2*1 = 1$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Original Series	1	0	2	3	5	1	2	3
Kernel		-1	2	1				
Convolution Series	1	7						

$$(0, 2, 3) \cdot (-1, 2, 1) = 0*(-1) + 2*2 + 3*1 = 7$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Original Series	1	0	2	3	5	1	2	3
Kernel			-1	2	1			
Convolution Series	1	7	9					

$$(2, 3, 5) \cdot (-1, 2, 1) = 2*(-1) + 3*2 + 5*1 = 9$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Original Series	1	0	2	3	5	1	2	3
Kernel				-1	2	1		
Convolution Series	1	7	9	8				

$$(3, 5, 1) \cdot (-1, 2, 1) = 3 * (-1) + 5 * 2 + 1 * 1 = 8$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Original Series	1	0	2	3	5	1	2	3
Kernel					-1	2	1	
Convolution Series	1	7	9	8	-1			

$$(5, 1, 2) \cdot (-1, 2, 1) = 5 * (-1) + 1 * 2 + 2 * 1 = -1$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Original Series	1	0	2	3	5	1	2	3
Kernel						-1	2	1
Convolution Series	1	7	9	8	-1	6		

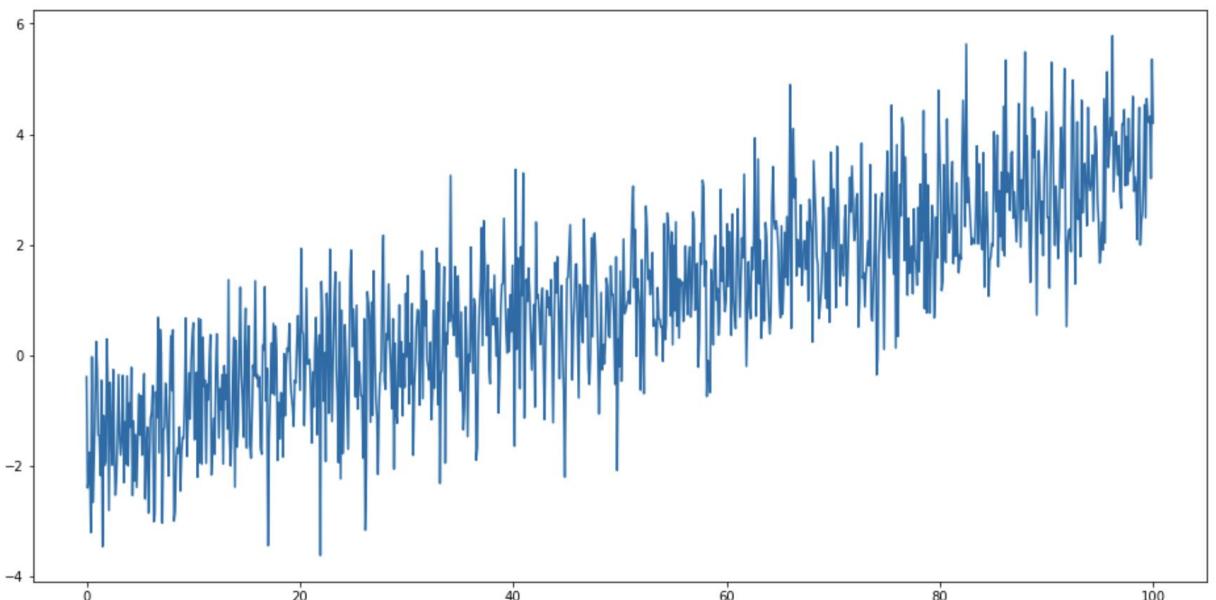
$$(1, 2, 3) \cdot (-1, 2, 1) = 1*(-1) + 2*2 + 3*1 = 6$$

# Convolution Operation (1D)

```
In [1]: import numpy as np  
import matplotlib.pyplot as plt
```

```
In [2]: N_series = 1000  
x = np.linspace(0, 100, N_series)  
y = 0.05 * x - 1.5 + np.random.randn(N_series)
```

```
In [3]: plt.figure(figsize=(16, 8))  
plt.plot(x, y)  
plt.show()
```



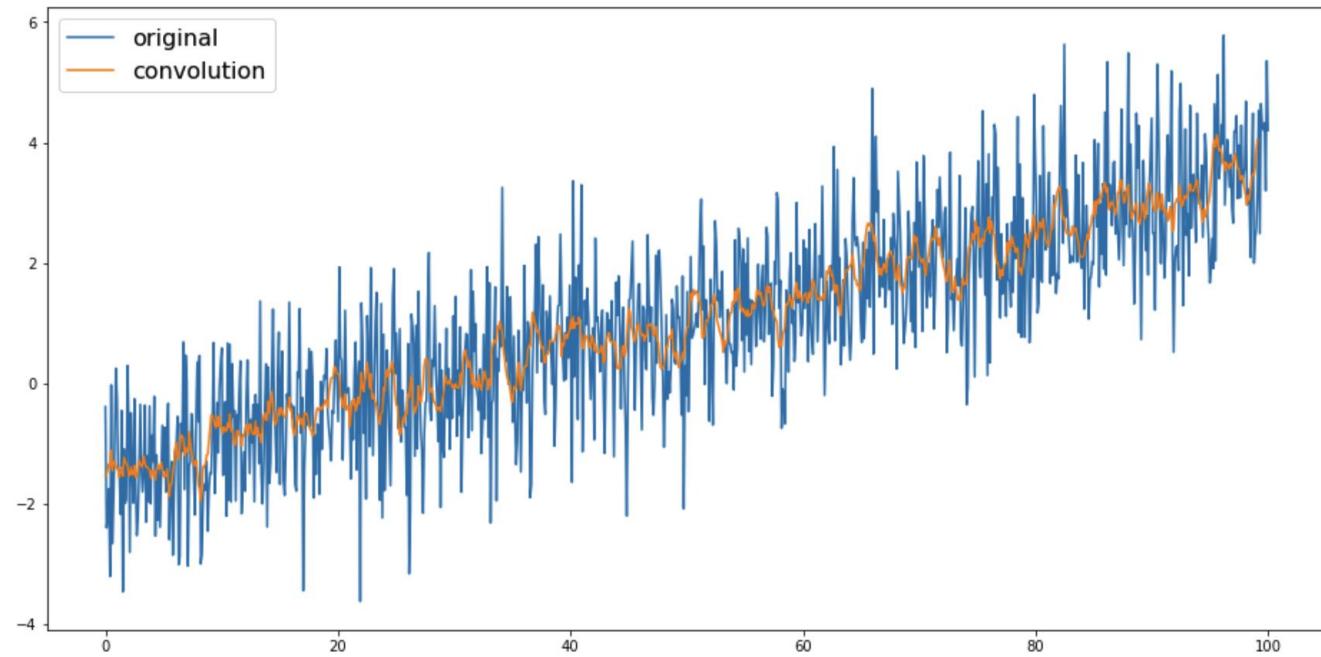
# Convolution Operation (1D)

```
In [4]: kernel = 0.1 * np.ones(10)
```

```
In [5]: def conv_1D(x, kernel):
    output = np.zeros((len(x) - len(kernel) + 1))
    for i in range(len(output)):
        output[i] = np.dot(kernel, x[i:i+len(kernel)])
    return output
```

```
In [6]: z = conv_1D(y, kernel)
```

```
In [7]: plt.figure(figsize=(16, 8))
plt.plot(x, y, label='original')
plt.plot(x[:len(z)], z, label='convolution')
plt.legend(fontsize=16)
plt.show()
```



# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Zero padding

Original Series	0	1	0	2	3	5	1	2	3	0
Kernel	-1	2	1							
Convolution Series		2								

$$(0, 1, 0) \cdot (-1, 2, 1) = 0*(-1) + 1*2 + 0*1 = 2$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Zero padding

Original Series	0	1	0	2	3	5	1	2	3	0
Kernel		-1	2	1						
Convolution Series		2	1							

$$(1, 0, 2) \cdot (-1, 2, 1) = 1*(-1) + 0*2 + 2*1 = 1$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Zero padding

Original Series		0	1	0	2	3	5	1	2	3	0
Kernel				-1	2	1					
Convolution Series		2	1	7							

$$(0, 2, 3) \cdot (-1, 2, 1) = 0*(-1) + 2*2 + 3*1 = 7$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Zero padding

Original Series	0	1	0	2	3	5	1	2	3	0
Kernel				-1	2	1				
Convolution Series	2	1	7	9						

$$(2, 3, 5) \cdot (-1, 2, 1) = 2 * (-1) + 3 * 2 + 5 * 1 = 9$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Zero padding

Original Series	0	1	0	2	3	5	1	2	3	0
Kernel					-1	2	1			
Convolution Series	2	1	7	9	8					

$$(3, 5, 1) \cdot (-1, 2, 1) = 3 * (-1) + 5 * 2 + 1 * 1 = 8$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 1

Zero padding

Original Series	0	1	0	2	3	5	1	2	3	0
Kernel						-1	2	1		
Convolution Series	2	1	7	9	8	-1				

$$(5, 1, 2) \cdot (-1, 2, 1) = 5 * (-1) + 1 * 2 + 2 * 1 = -1$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series sliding with stride of 1

Zero padding

Original Series	0	1	0	2	3	5	1	2	3	0
Kernel						-1	2	1		
Convolution Series	2	1	7	9	8	-1	6			

$$(1, 2, 3) \cdot (-1, 2, 1) = 1 * (-1) + 2 * 2 + 3 * 1 = 6$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series sliding with stride of 1

Zero padding

Original Series	0	1	0	2	3	5	1	2	3	0
Kernel							-1	2	1	
Convolution Series	2	1	7	9	8	-1	6	4		

$$(2, 3, 0) \cdot (-1, 2, 1) = 2 * (-1) + 3 * 2 + 0 * 1 = 4$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 2

Zero padding

Original Series			0	1	0	2	3	5	1	2	0
Kernel			-1	2	1						
Convolution Series			2								

$$(0, 1, 0) \cdot (-1, 2, 1) = 0*(-1) + 1*2 + 0*1 = 2$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 2

Zero padding

Original Series		0	1	0	2	3	5	1	2	0
Kernel				-1	2	1				
Convolution Series		2	7							

$$(0, 2, 3) \cdot (-1, 2, 1) = 0*(-1) + 2*2 + 3*1 = 7$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 2

Zero padding

Original Series		0	1	0	2	3	5	1	2	0
Kernel					-1	2	1			
Convolution Series		2	7	8						

$$(3, 5, 1) \cdot (-1, 2, 1) = 3 * (-1) + 5 * 2 + 1 * 1 = 8$$

# Convolution Operation (1D)

Discrete 1-dim convolution with a kernel of size 3

Dot product of the kernel and the corresponding part of the original series  
sliding with stride of 2

Zero padding

Original Series	0	1	0	2	3	5	1	2	0
Kernel						-1	2	1	
Convolution Series	2	7	8	3					

$$(1, 2, 0) \cdot (-1, 2, 1) = 1 * (-1) + 2 * 2 + 0 * 1 = 3$$

# Convolution (1D) Summing Up

Input of size W

Kernel of size K

Stride of S

Padding of P

Output size = ?



# Convolution (1D) Summing Up

Input of size W

Padding of P

Kernel of size K

Stride of S

$$\text{Output size} = (W + 2*P - K) / S + 1$$

=> Must be an integer for valid convolution!

# Convolution Operation (2D)

Discrete 2-dim convolution with a kernel of size 3x3

Dot product of the kernel and the corresponding part of the input matrix  
sliding with stride of 1

Input Matrix						Kernel			Output Matrix			
1	0	2	1	-1	1	-1	0	1				
0	1	-1	1	0	2	0	2	1				
1	0	0	0	-1	1	1	0	-1				
0	-1	-1	0	2	0							
1	1	0	-2	1	0							
0	0	1	-1	0	1							

$$(1, 0, 2, 0, 1, -1, 1, 0, 0) \cdot (-1, 0, 1, 0, 2, 1, 1, 0, -1) = 3$$

# Convolution Operation (2D)

Discrete 2-dim convolution with a kernel of size 3x3

Dot product of the kernel and the corresponding part of the input matrix  
sliding with stride of 1

Input Matrix						Kernel				Output Matrix			
1	0	2	1	-1	1		-1	0	1				
0	1	-1	1	0	2		0	2	1				
1	0	0	0	-1	1		1	0	-1				
0	-1	-1	0	2	0								
1	1	0	-2	1	0								
0	0	1	-1	0	1								

$$(0, 2, 1, 1, -1, 1, 0, 0, 0) \cdot (-1, 0, 1, 0, 2, 1, 1, 0, -1) = 0$$

# Convolution Operation (2D)

Discrete 2-dim convolution with a kernel of size 3x3

Dot product of the kernel and the corresponding part of the input matrix  
sliding with stride of 1

Input Matrix						Kernel				Output Matrix			
1	0	2	1	-1	1			-1	0	1			
0	1	-1	1	0	2			0	2	1			
1	0	0	0	-1	1			1	0	-1			
0	-1	-1	0	2	0								
1	1	0	-2	1	0								
0	0	1	-1	0	1								

$$(2, 1, -1, -1, 1, 0, 0, 0, -1) \cdot (-1, 0, 1, 0, 2, 1, 1, 0, -1) = 0$$

# Convolution Operation (2D)

Discrete 2-dim convolution with a kernel of size 3x3

Dot product of the kernel and the corresponding part of the input matrix  
sliding with stride of 1

Input Matrix						Kernel			Output Matrix			
1	0	2	1	-1	1				-1	0	1	
0	1	-1	1	0	2				0	2	1	
1	0	0	0	-1	1				1	0	-1	
0	-1	-1	0	2	0							
1	1	0	-2	1	0							
0	0	1	-1	0	1							

$$(1, -1, 1, 1, 0, 2, 0, -1, 1) \cdot (-1, 0, 1, 0, 2, 1, 1, 0, -1) = 2$$

# Convolution Operation (2D)

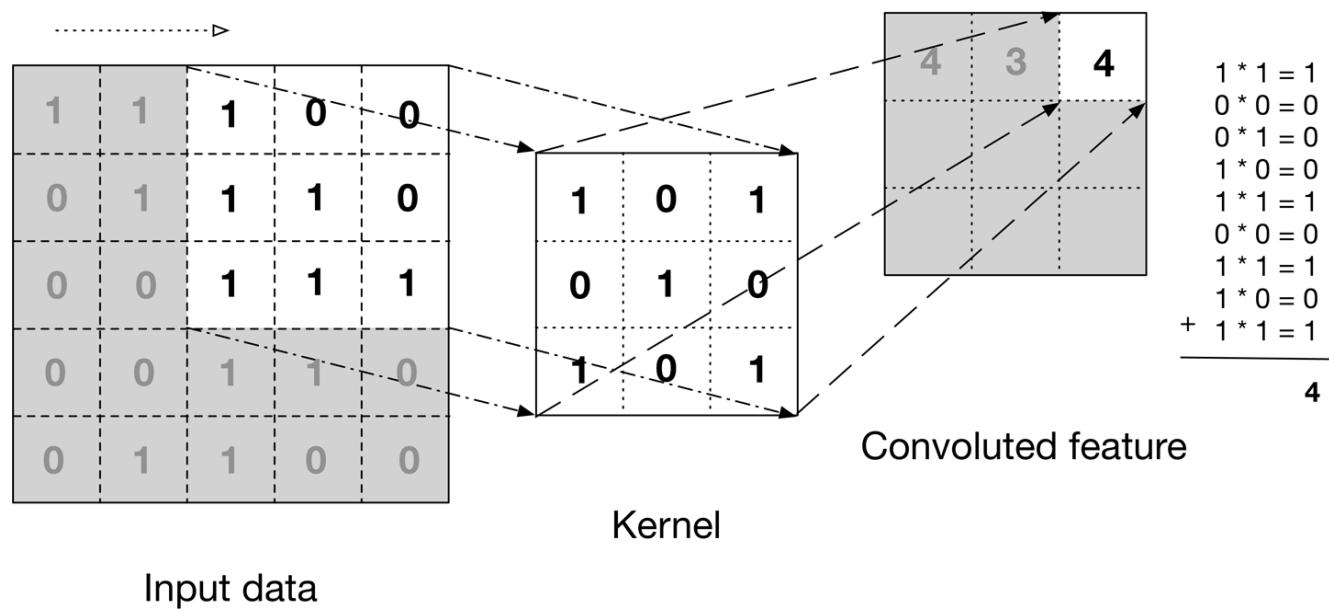
Discrete 2-dim convolution with a kernel of size 3x3

Dot product of the kernel and the corresponding part of the input matrix  
sliding with stride of 1

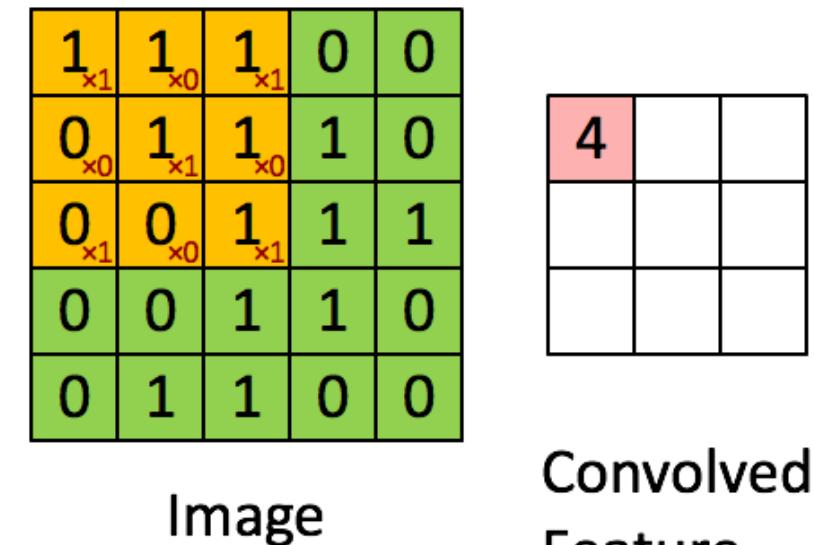
Input Matrix						Kernel						Output Matrix			
1	0	2	1	-1	1							3	0	0	2
0	1	-1	1	0	2		-1	0	1			0			
1	0	0	0	-1	1		0	2	1						
0	-1	-1	0	2	0		1	0	-1						
1	1	0	-2	1	0										
0	0	1	-1	0	1										

$$(0, 1, -1, 1, 0, 0, 0, -1, -1) \cdot (-1, 0, 1, 0, 2, 1, 1, 0, -1) = 0$$

# Convolution Operation (2D)



$$\begin{array}{r}
 & 1 * 1 = 1 \\
 & 0 * 0 = 0 \\
 & 0 * 1 = 0 \\
 & 1 * 0 = 0 \\
 & 1 * 1 = 1 \\
 & 0 * 0 = 0 \\
 & 1 * 1 = 1 \\
 & 1 * 0 = 0 \\
 + & 1 * 1 = 1 \\
 \hline
 & 4
 \end{array}$$

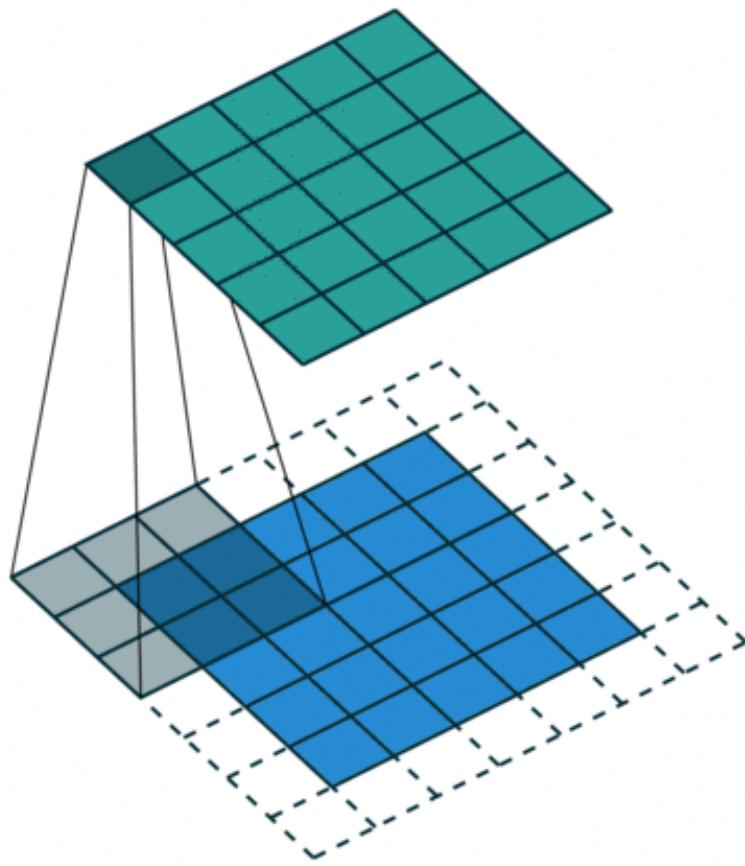


## Convolved Feature

<https://www.analyticsvidhya.com/blog/2021/05/convolutional-neural-networks-cnn/>

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a5>

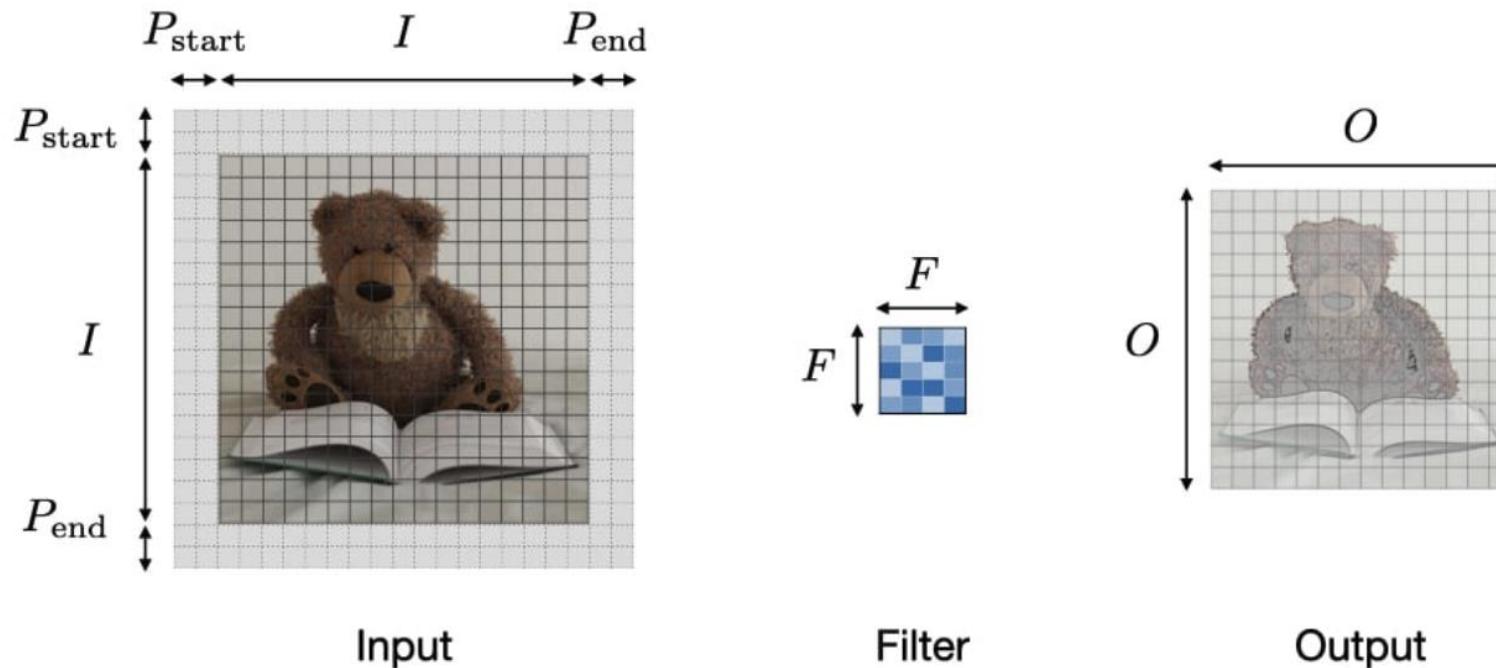
# Padding in 2D Convolution



[https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

# Input and Output Sizes

$$O = \frac{I - F + P_{\text{start}} + P_{\text{end}}}{S} + 1$$



# Effects of Convolution Kernels (Filters)

 $*$ 

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

 $=$ 

Original

Identical image

 $*$ 

$$\begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

 $=$ Shifted left  
By 1 pixel $*$ 

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

 $=$ 

Original

Blur (with a mean filter)

 $*$ 

$$\left( \begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \right)$$

$$\begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

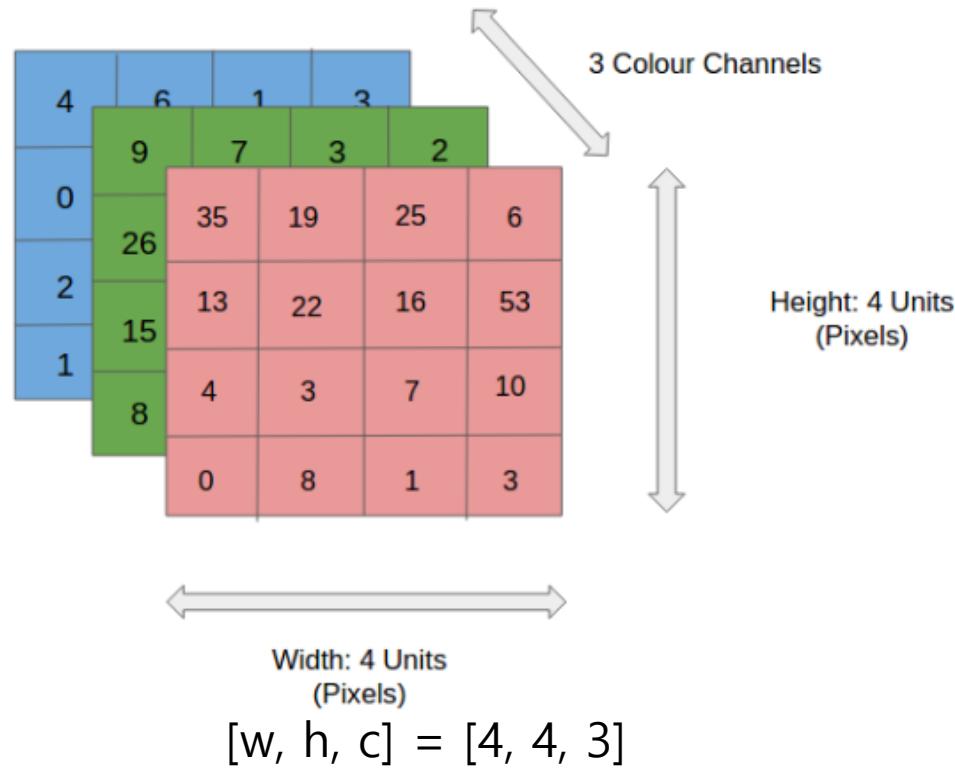
$$\begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Sharpening filter  
(accentuates edges)

# Color Image has Three Channels! (RGB)

Color input image = Tensor of rank 3

Tensor of rank N = N-dim array of numbers (simplified definition)



<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

# Kernel has parameters for all input channels (depth)

Kernel size: [w, h, c] = [3, 3, 3]

Note: This is still 2D convolution! (Why?)

0	0	0	0	0	0	0	...
0	156	155	156	158	158	158	...
0	153	154	157	159	159	159	...
0	149	151	155	158	159	159	...
0	146	146	149	153	158	158	...
0	145	143	143	148	158	158	...
...	...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	169	...
0	164	165	168	170	170	170	...
0	160	162	166	169	170	170	...
0	156	156	159	163	168	168	...
0	155	153	153	158	168	168	...
...	...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	165	...
0	160	161	164	166	166	166	...
0	156	158	162	165	166	166	...
0	155	155	158	162	167	167	...
0	154	152	152	157	167	167	...
...	...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

+

-498

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164

+

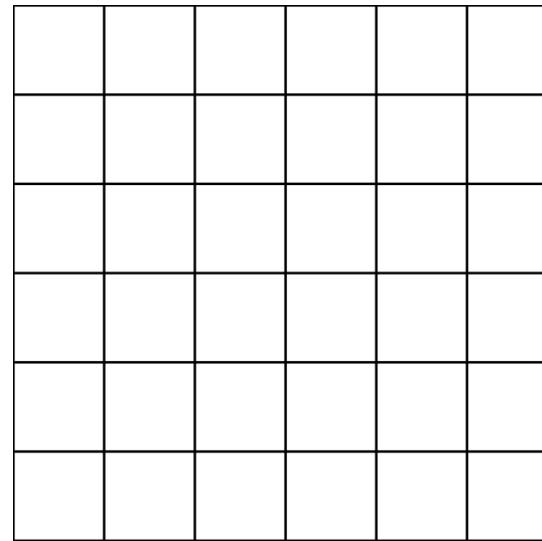
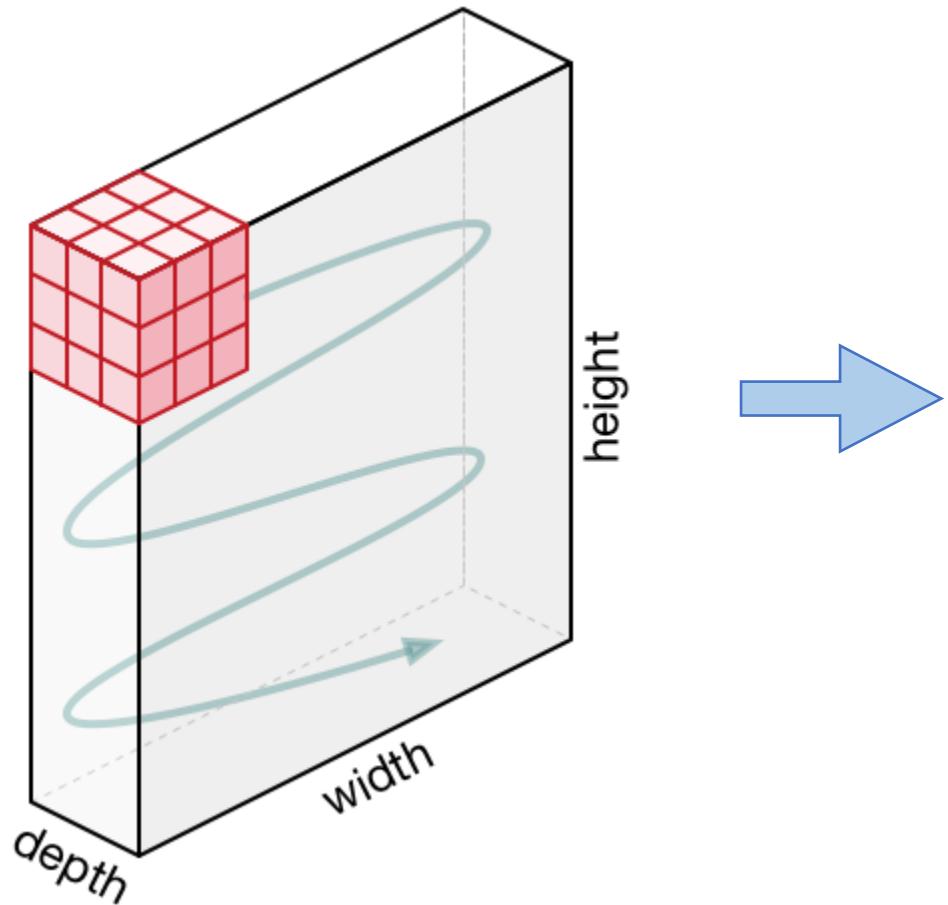
-25							...
							...
							...
							...
...	...	...	...	...	...	...	...

Output

Bias = 1

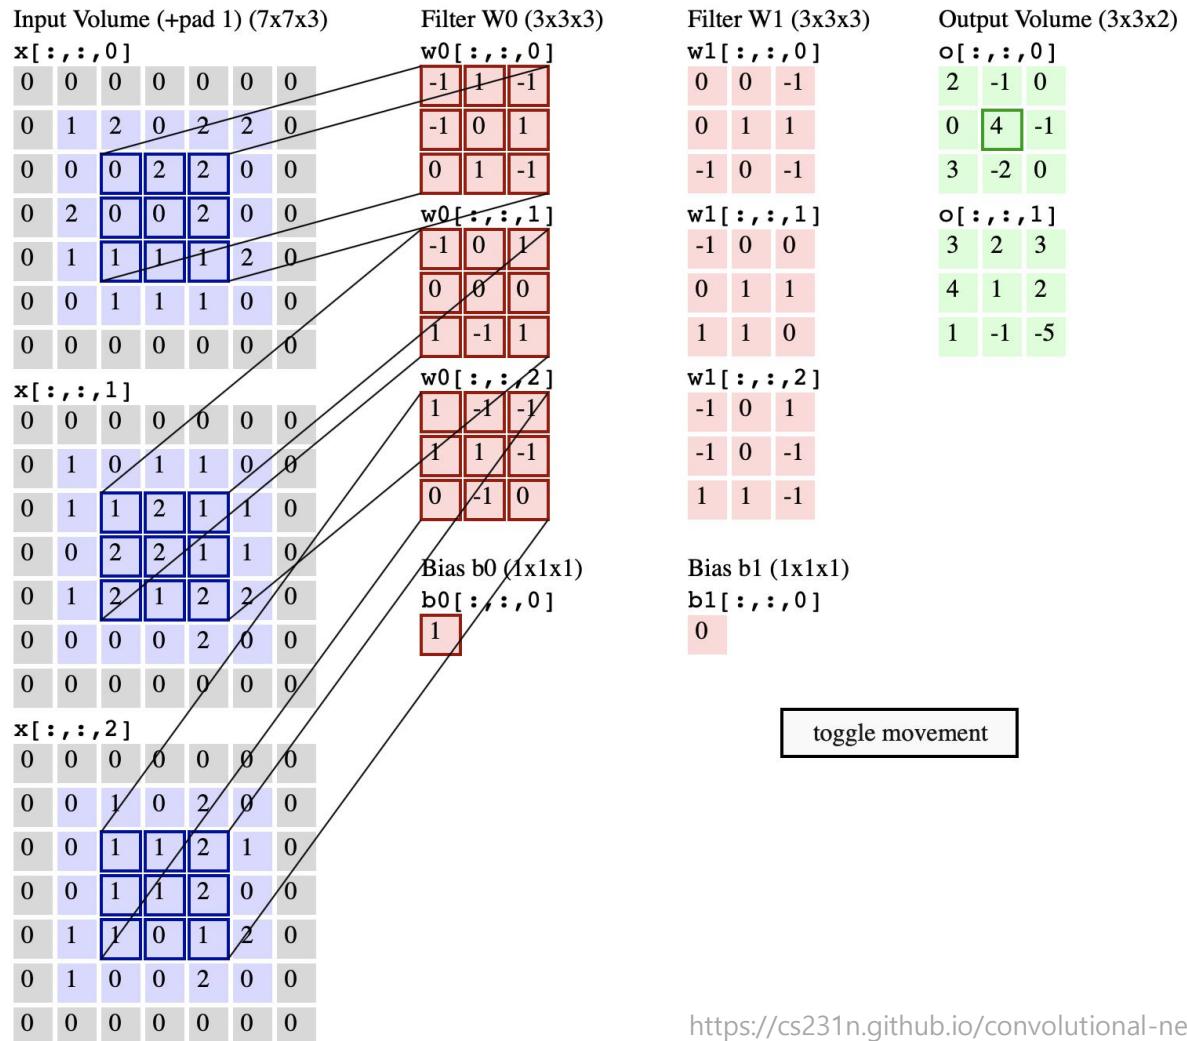
<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

## 2D Convolution may have a depth dim., but moves in 2D



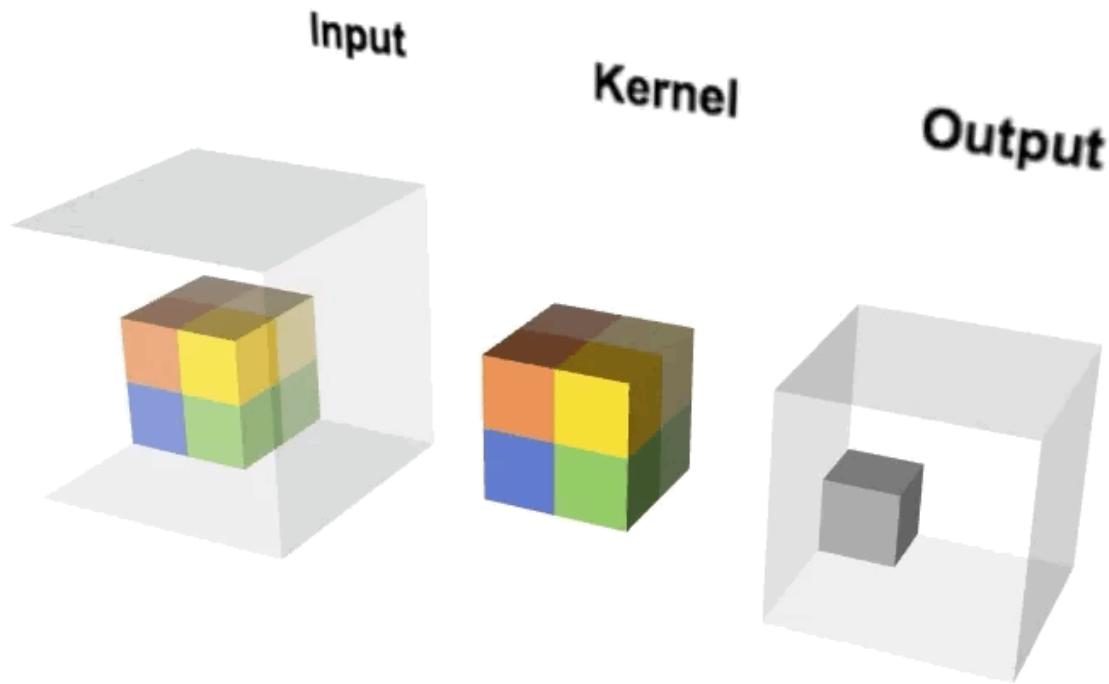
Output size: [w, h, 1]

# Multiple kernels (filters) for the same input image



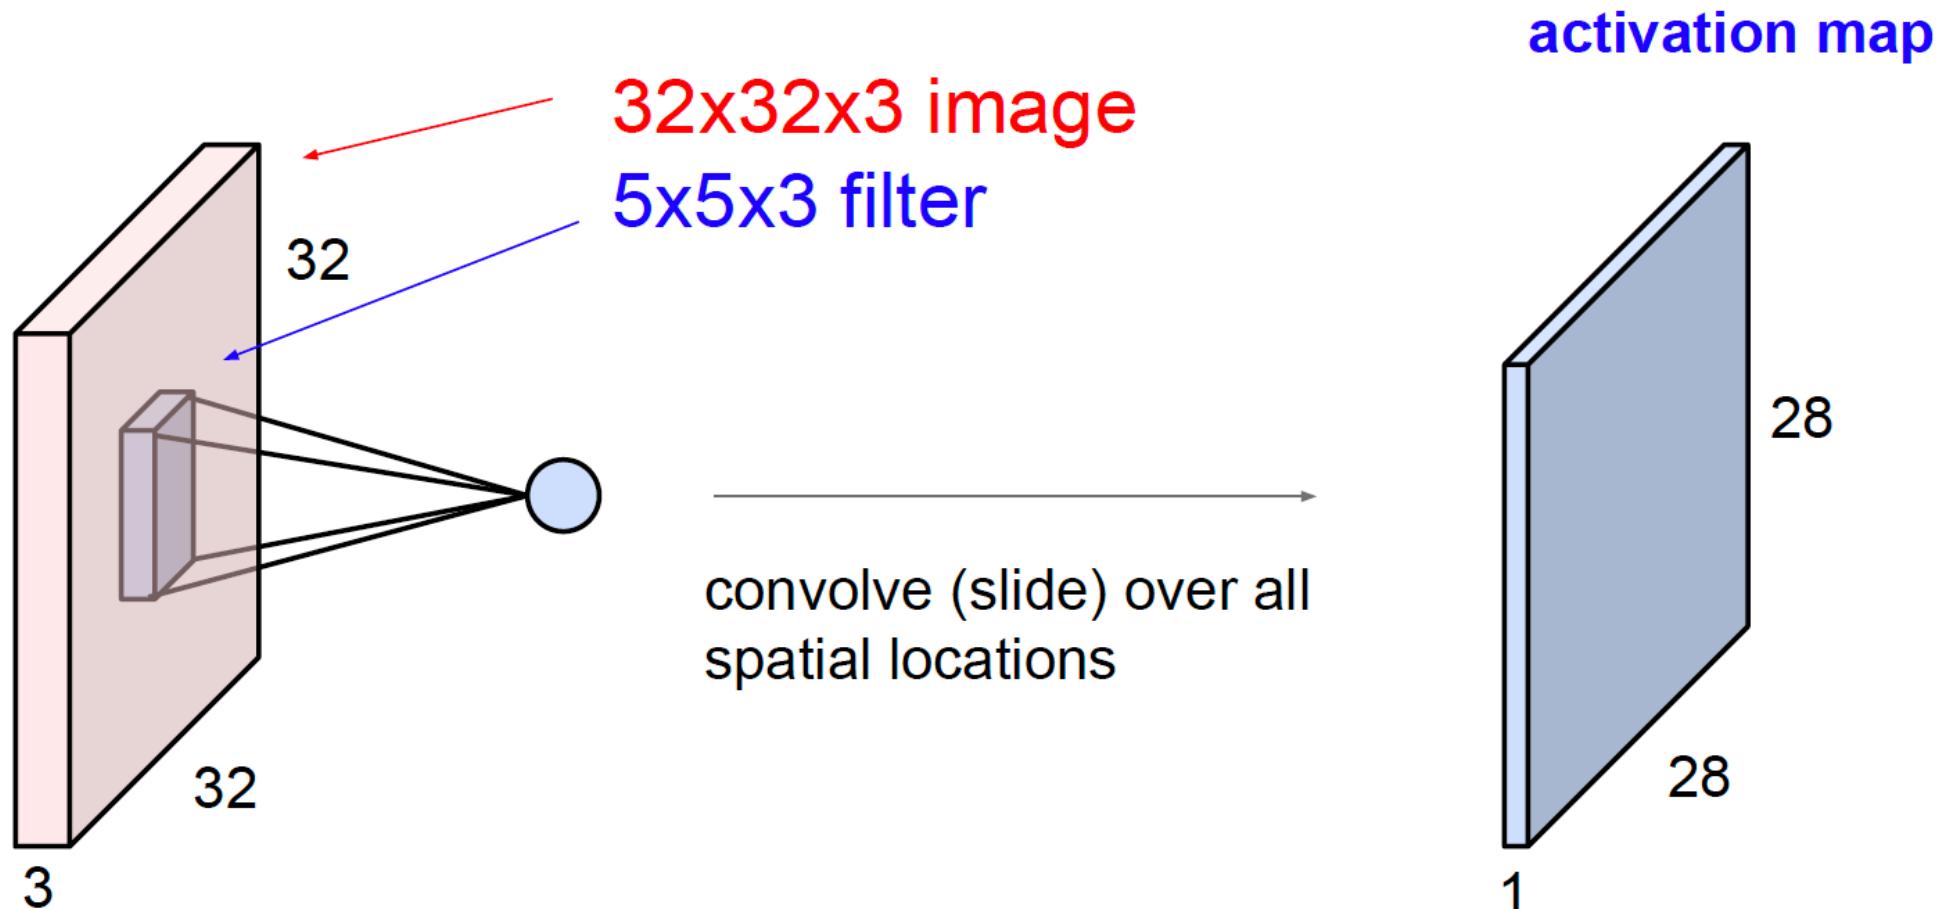
<https://cs231n.github.io/convolutional-networks/>

# 3D Convolution

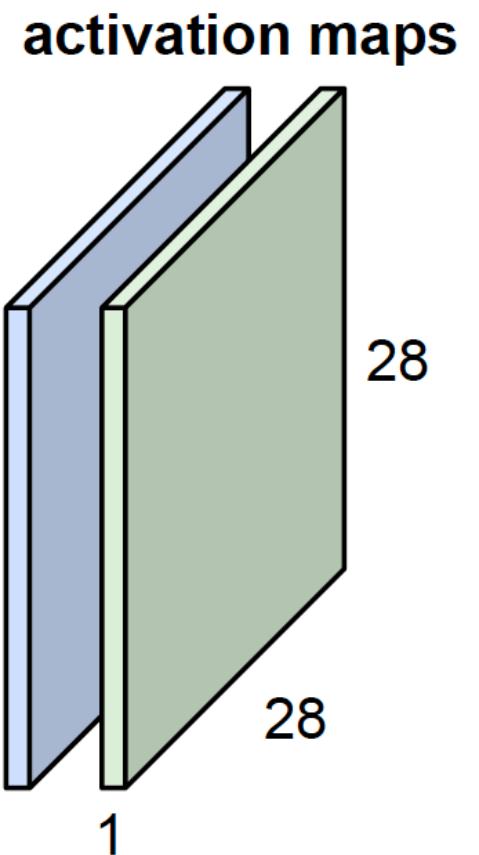
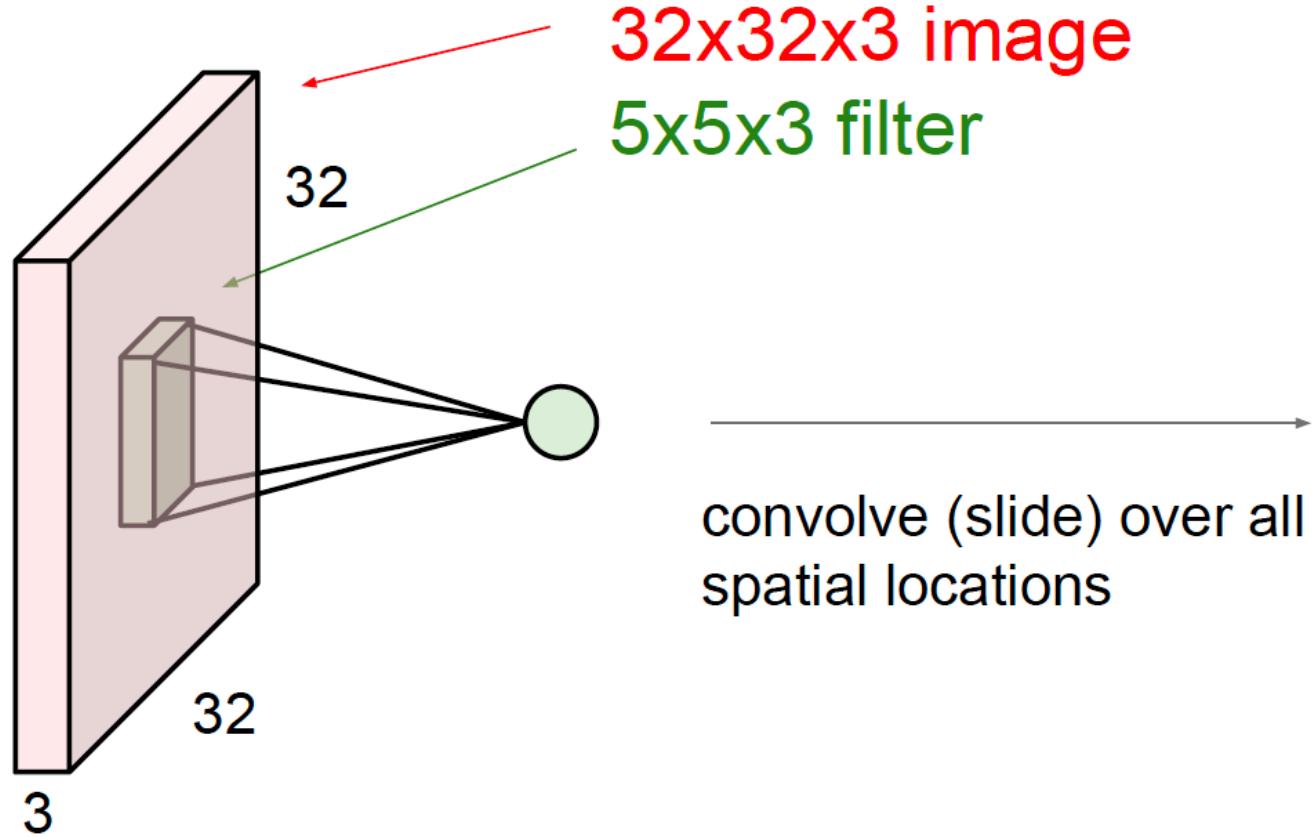


<https://thomelane.github.io/convolutions/3DConv.html>

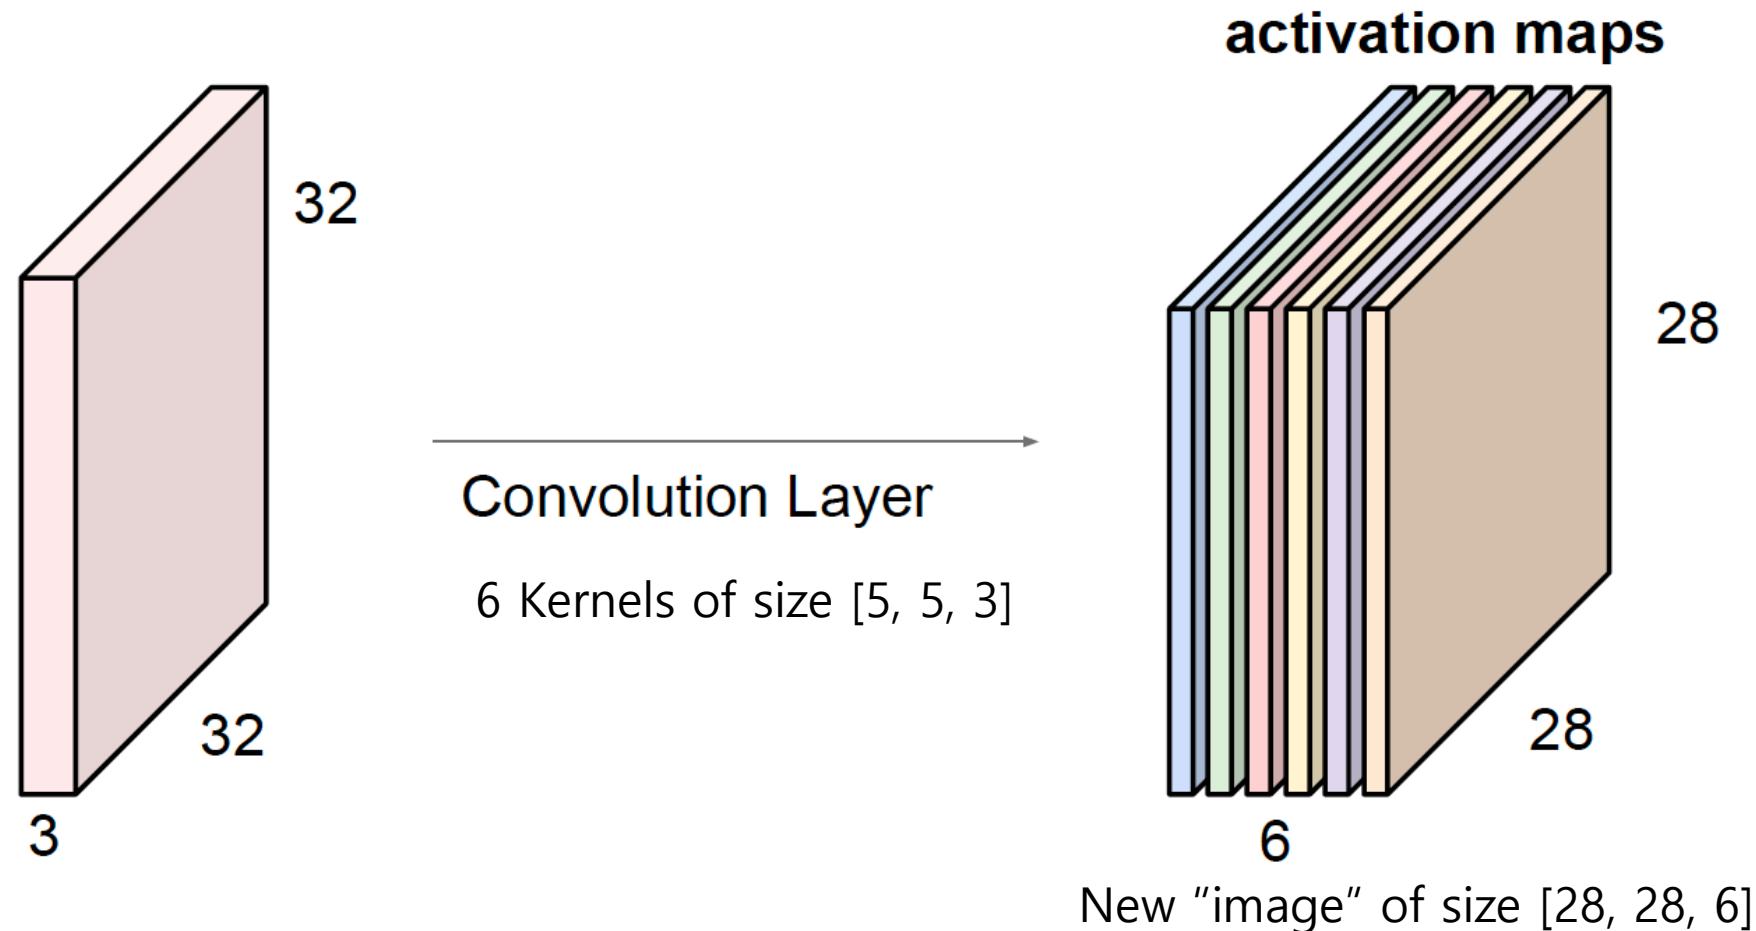
# Convolutional Layer



# Convolutional Layer

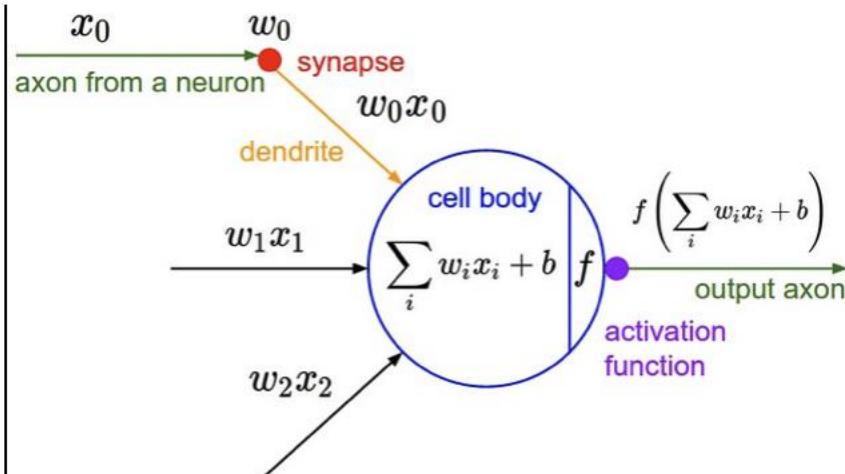
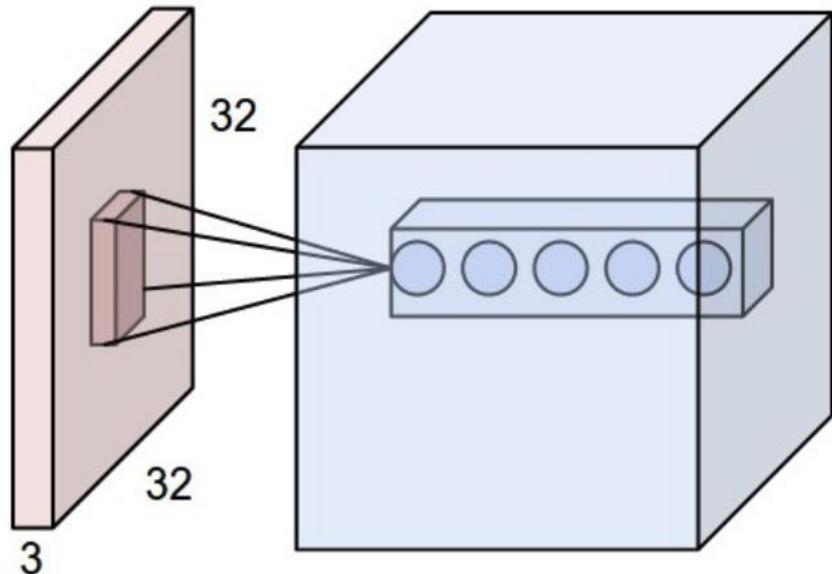


# Convolutional Layer



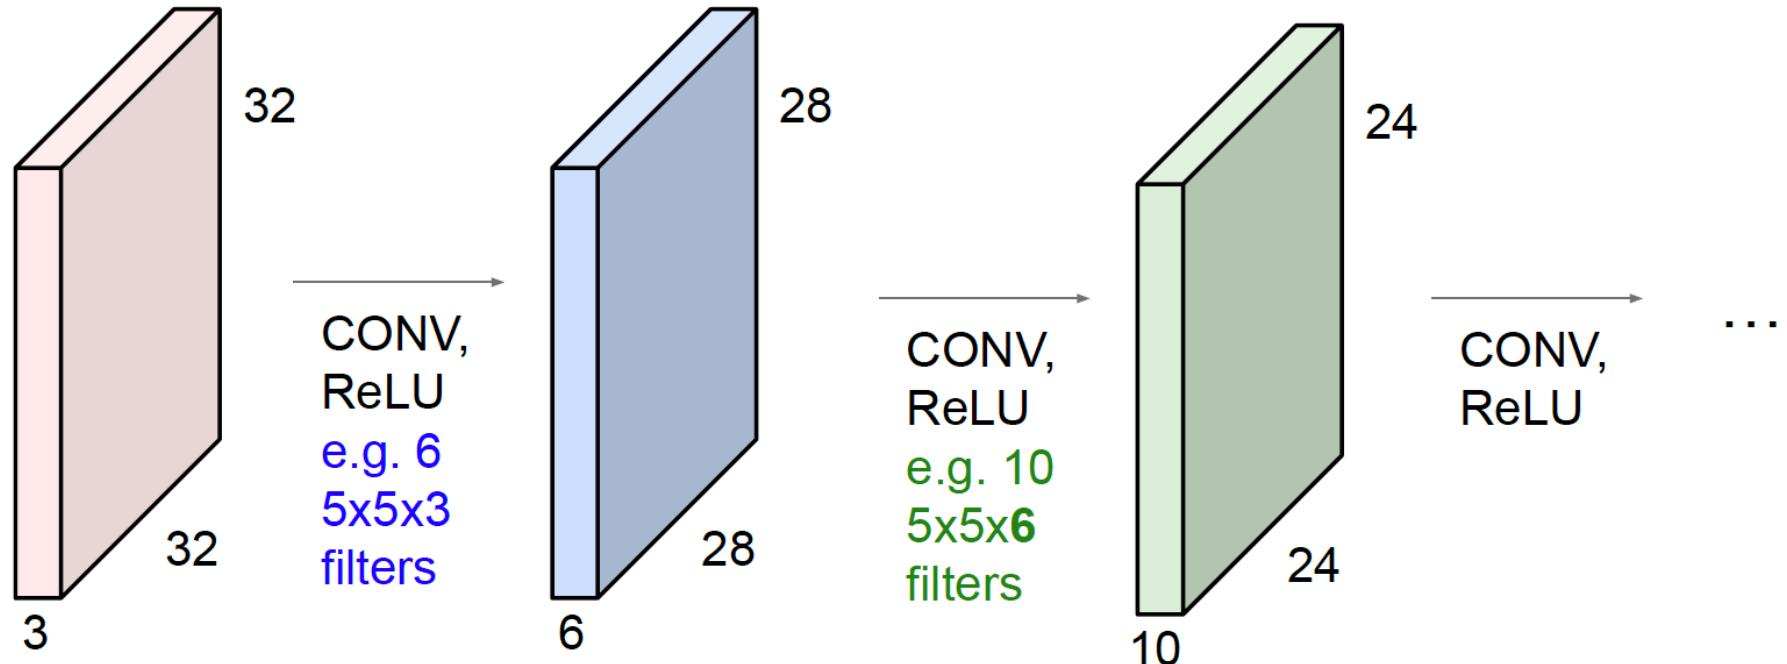
# Local Connectivity & Parameter Sharing

Convolution -> Local connectivity and parameter sharing of neurons



<https://cs231n.github.io/convolutional-networks/>

# Convolutional Layers

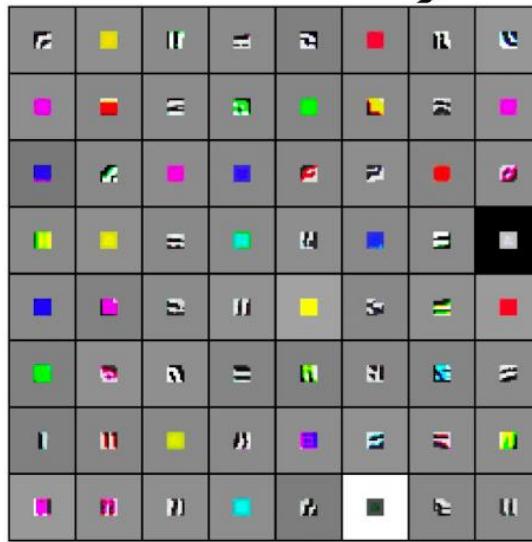


# Hierarchical Feature Learning

## Preview

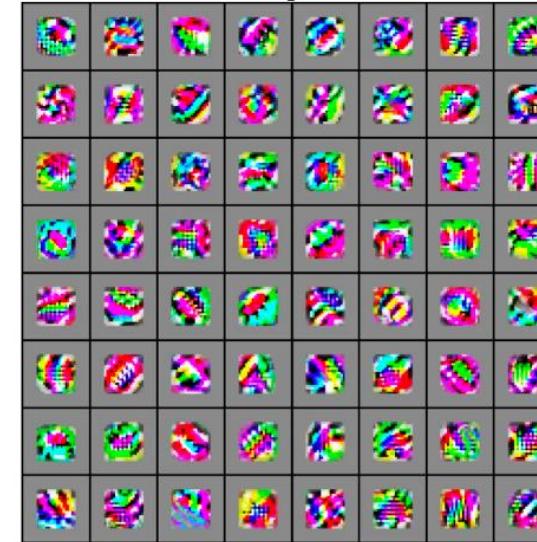


Low-level features

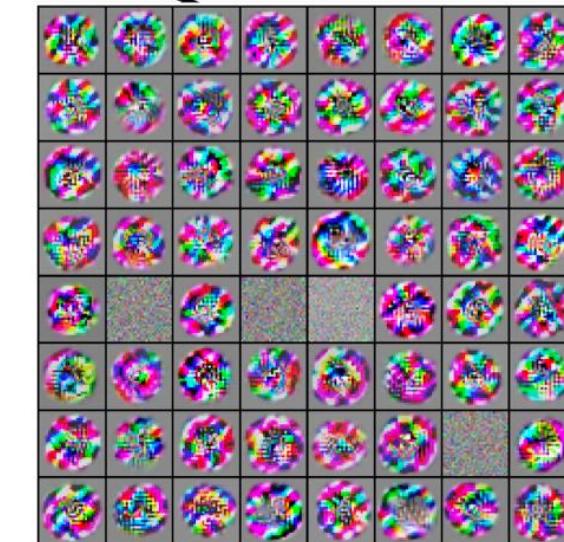


[Zeiler and Fergus 2013]

Mid-level features

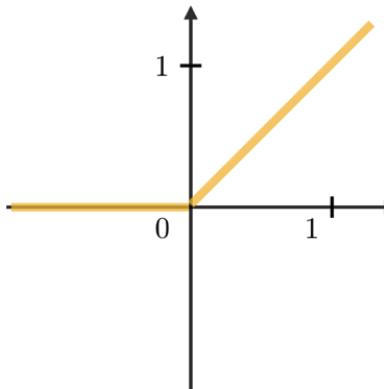
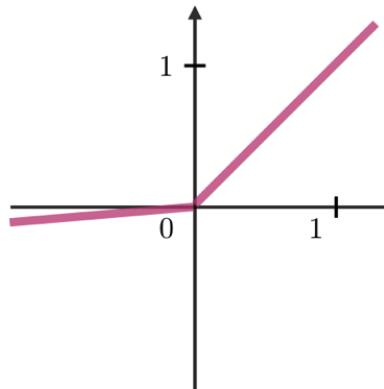
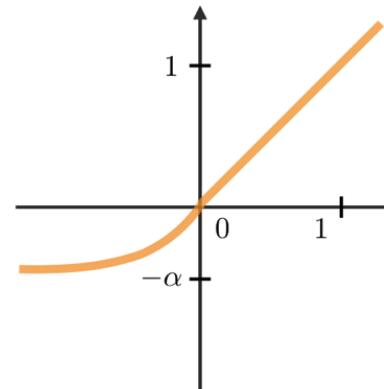


High-level features

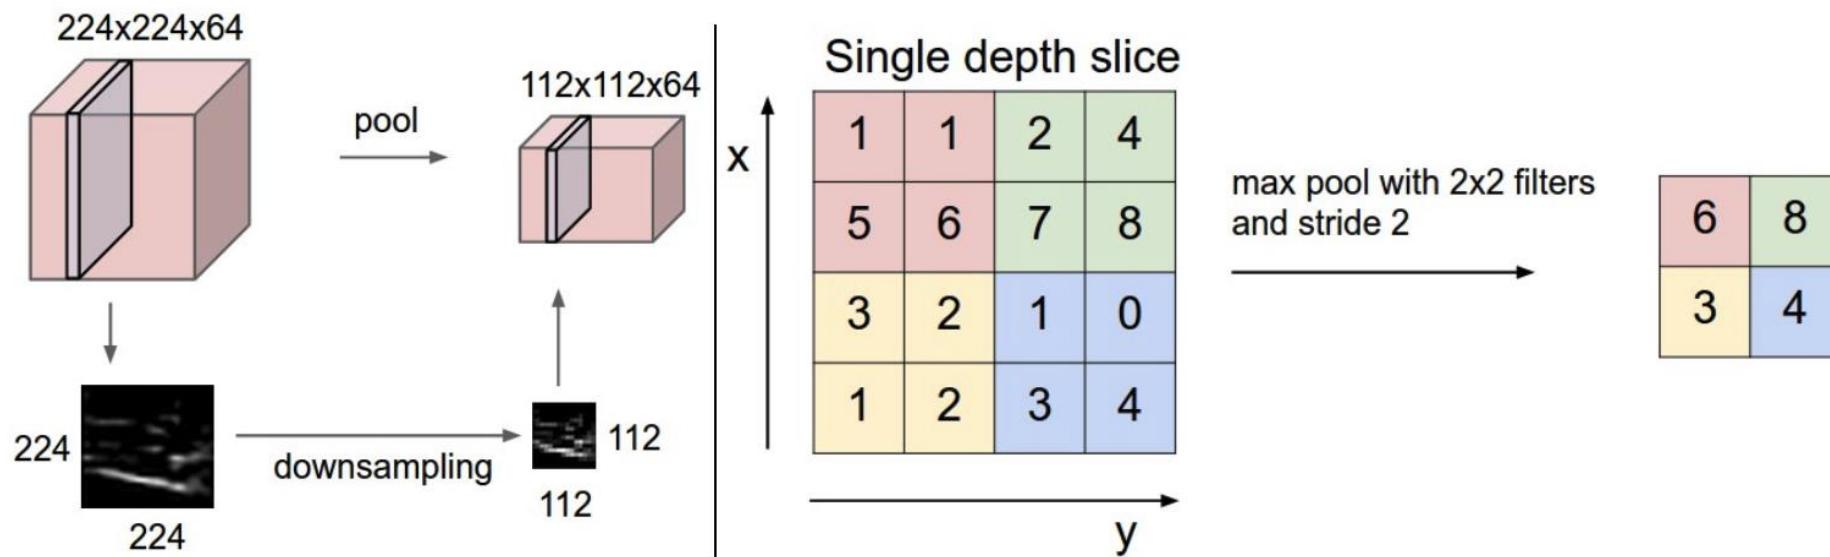


Visualization of VGG-16 by Lane McIntosh. VGG-16 architecture from [Simonyan and Zisserman 2014].

# Activation Functions for Convolutional Layers

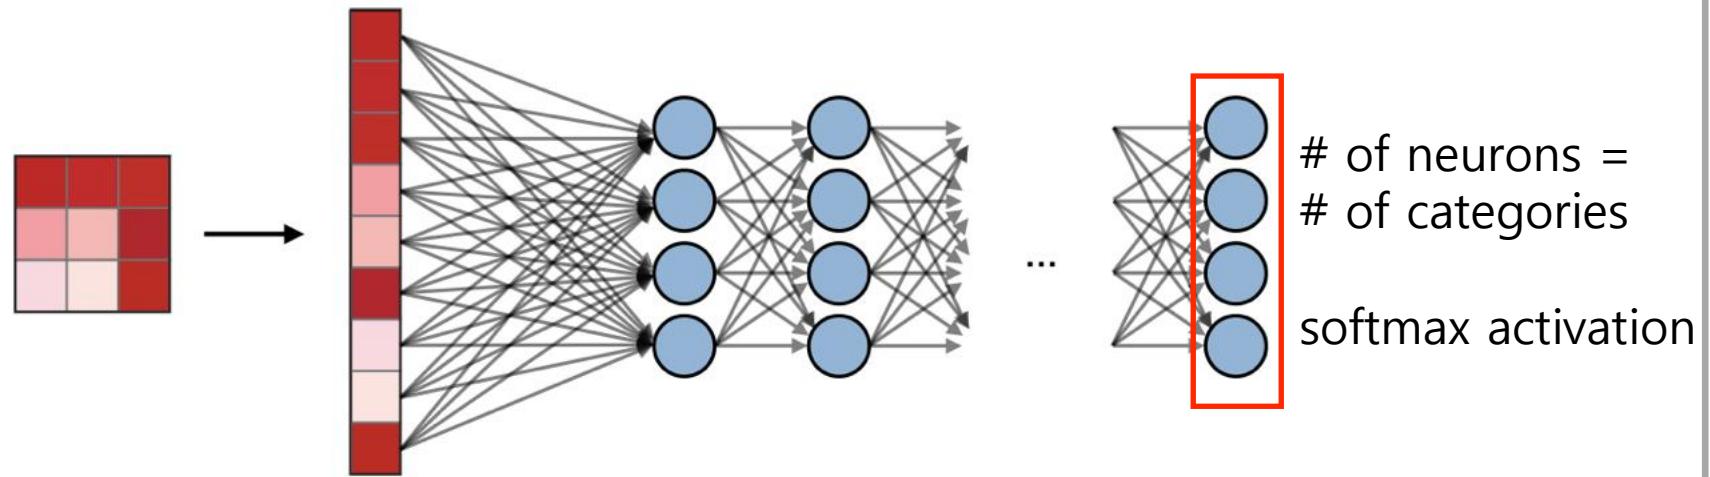
ReLU	Leaky ReLU	ELU
$g(z) = \max(0, z)$	$g(z) = \max(\epsilon z, z)$ with $\epsilon \ll 1$	$g(z) = \max(\alpha(e^z - 1), z)$ with $\alpha \ll 1$
		
<ul style="list-style-type: none"><li>• Non-linearity complexities biologically interpretable</li></ul>	<ul style="list-style-type: none"><li>• Addresses dying ReLU issue for negative values</li></ul>	<ul style="list-style-type: none"><li>• Differentiable everywhere</li></ul>

# Pooling Layer (Max or Average)

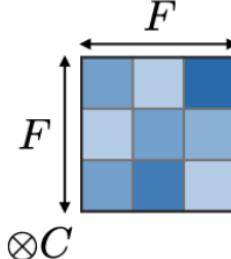
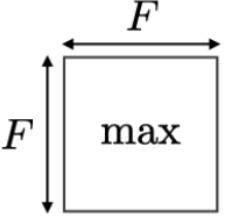
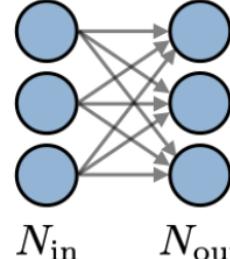


<https://cs231n.github.io/convolutional-networks/>

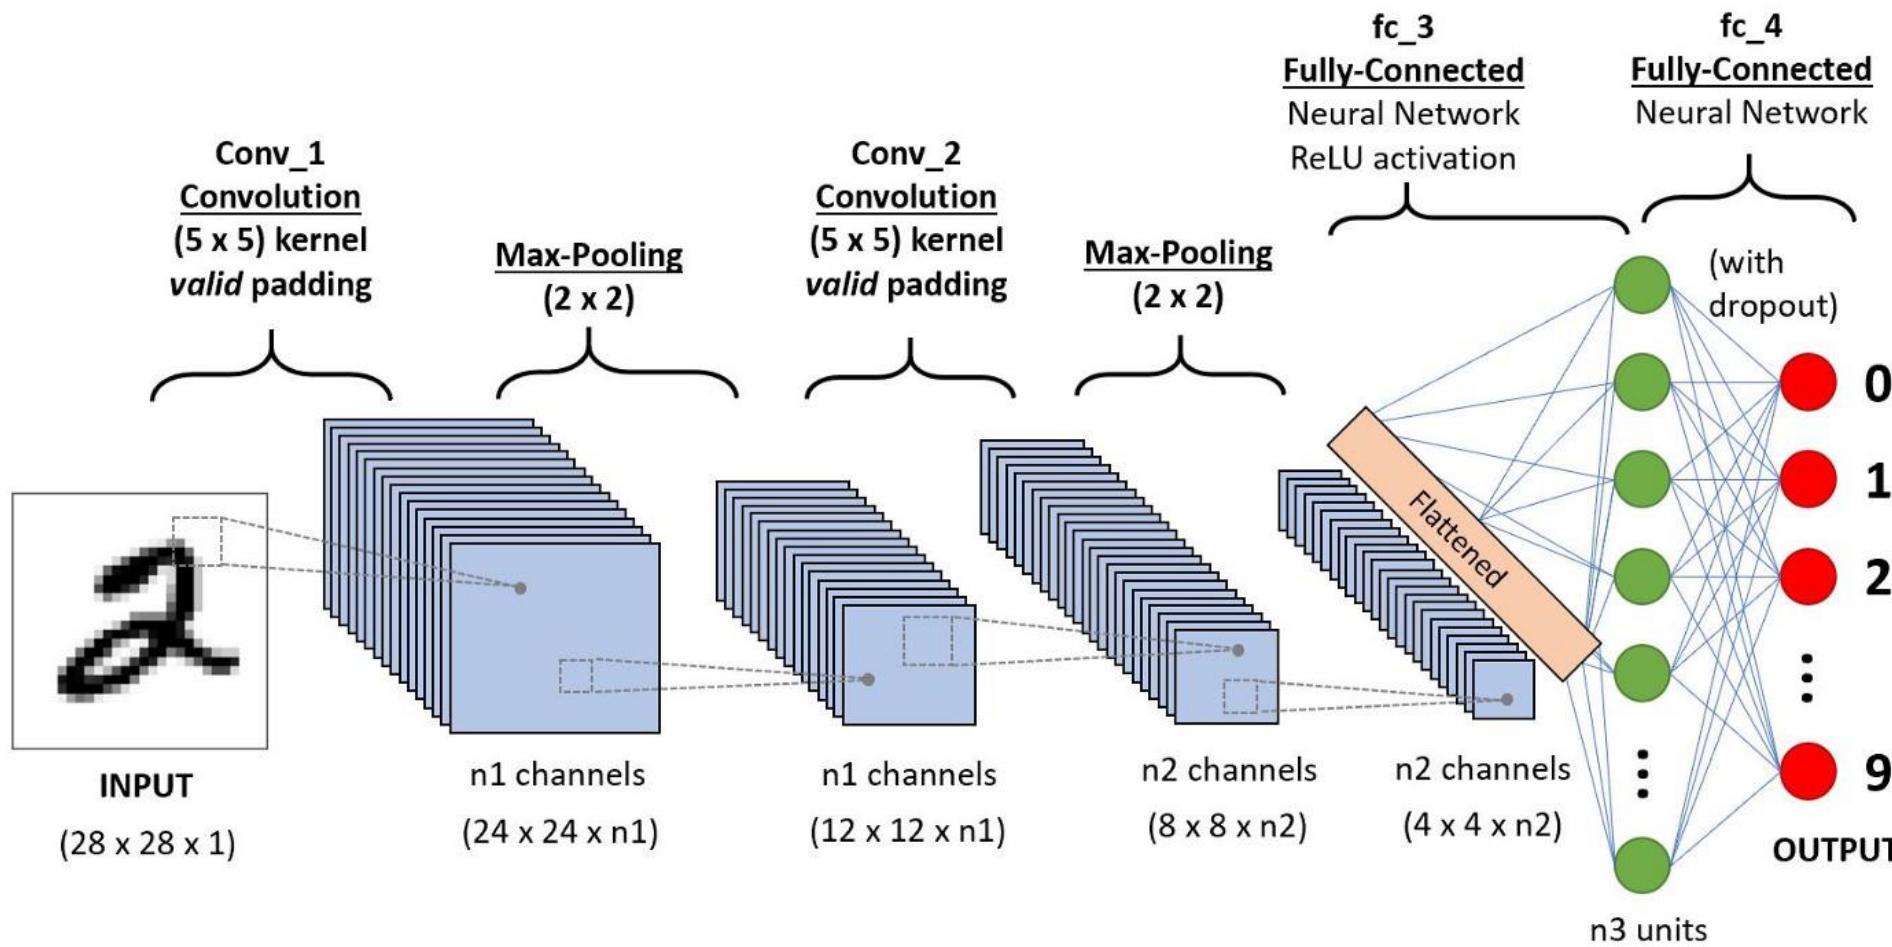
# Fully-connected Layer



# Summary for Layer Sizes

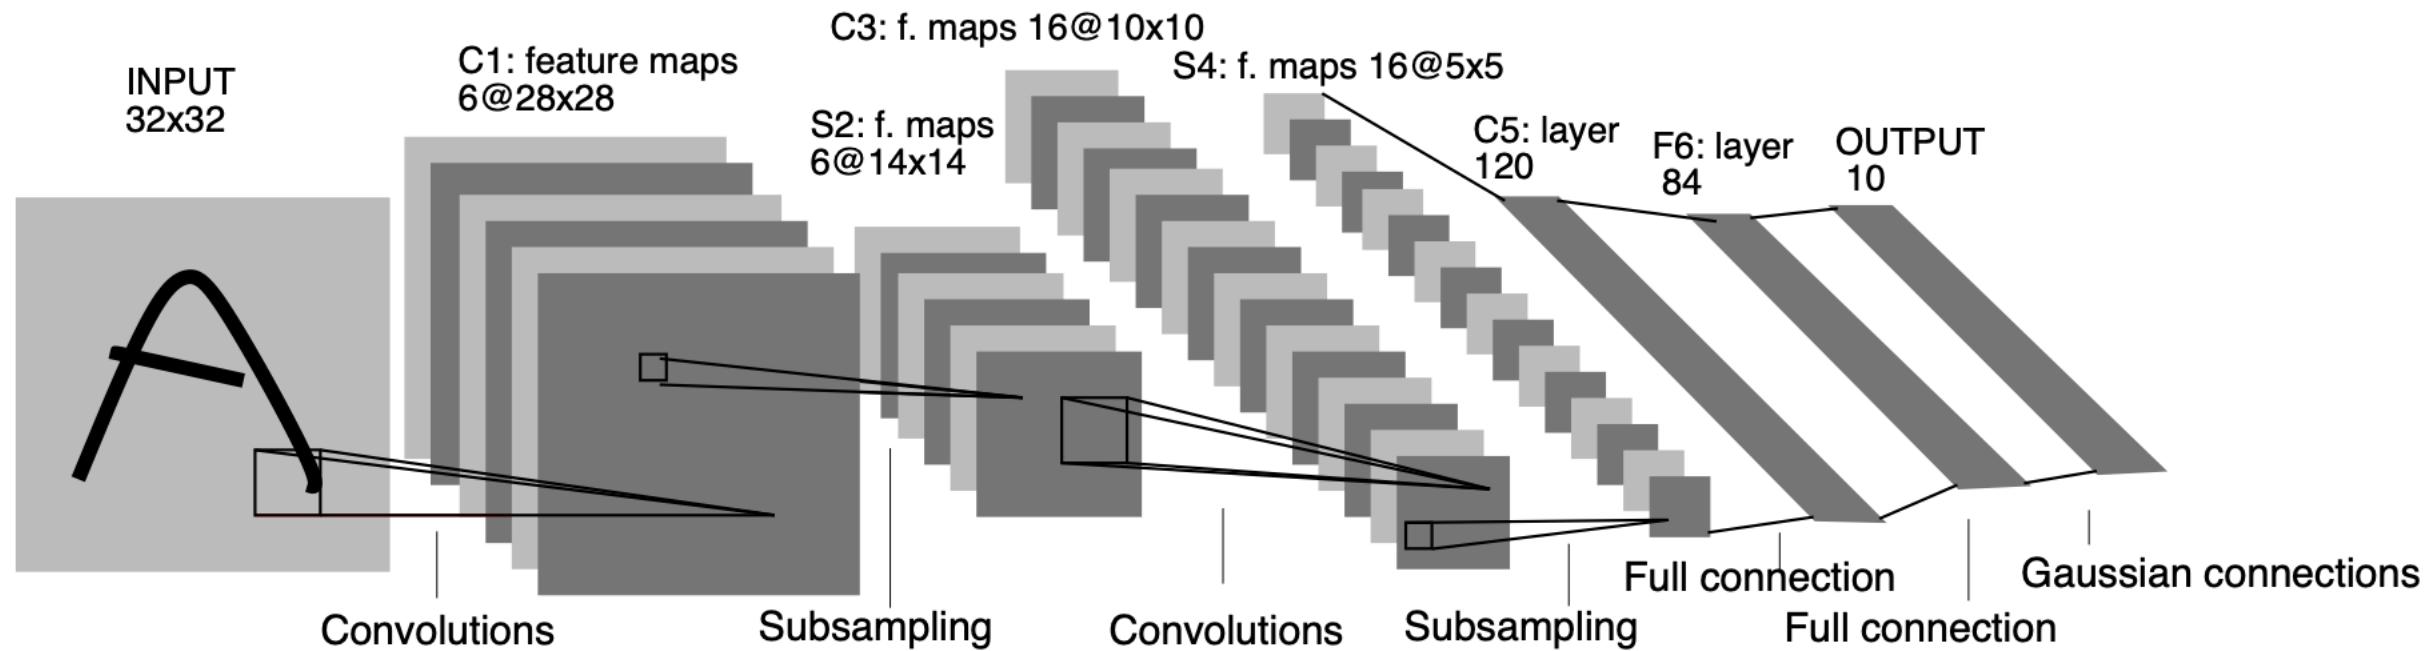
	<b>CONV</b>	<b>POOL</b>	<b>FC</b>
<b>Illustration</b>			
<b>Input size</b>	$I \times I \times C$	$I \times I \times C$	$N_{\text{in}}$
<b>Output size</b>	$O \times O \times K$	$O \times O \times C$	$N_{\text{out}}$
<b>Number of parameters</b>	$(F \times F \times C + 1) \cdot K$	0	$(N_{\text{in}} + 1) \times N_{\text{out}}$
<b>Remarks</b>	<ul style="list-style-type: none"> <li>One bias parameter per filter</li> <li>In most cases, <math>S &lt; F</math></li> <li>A common choice for <math>K</math> is <math>2C</math></li> </ul>	<ul style="list-style-type: none"> <li>Pooling operation done channel-wise</li> <li>In most cases, <math>S = F</math></li> </ul>	<ul style="list-style-type: none"> <li>Input is flattened</li> <li>One bias parameter per neuron</li> <li>The number of FC neurons is free of structural constraints</li> </ul>

# Example of a Simple Convolutional Neural Network



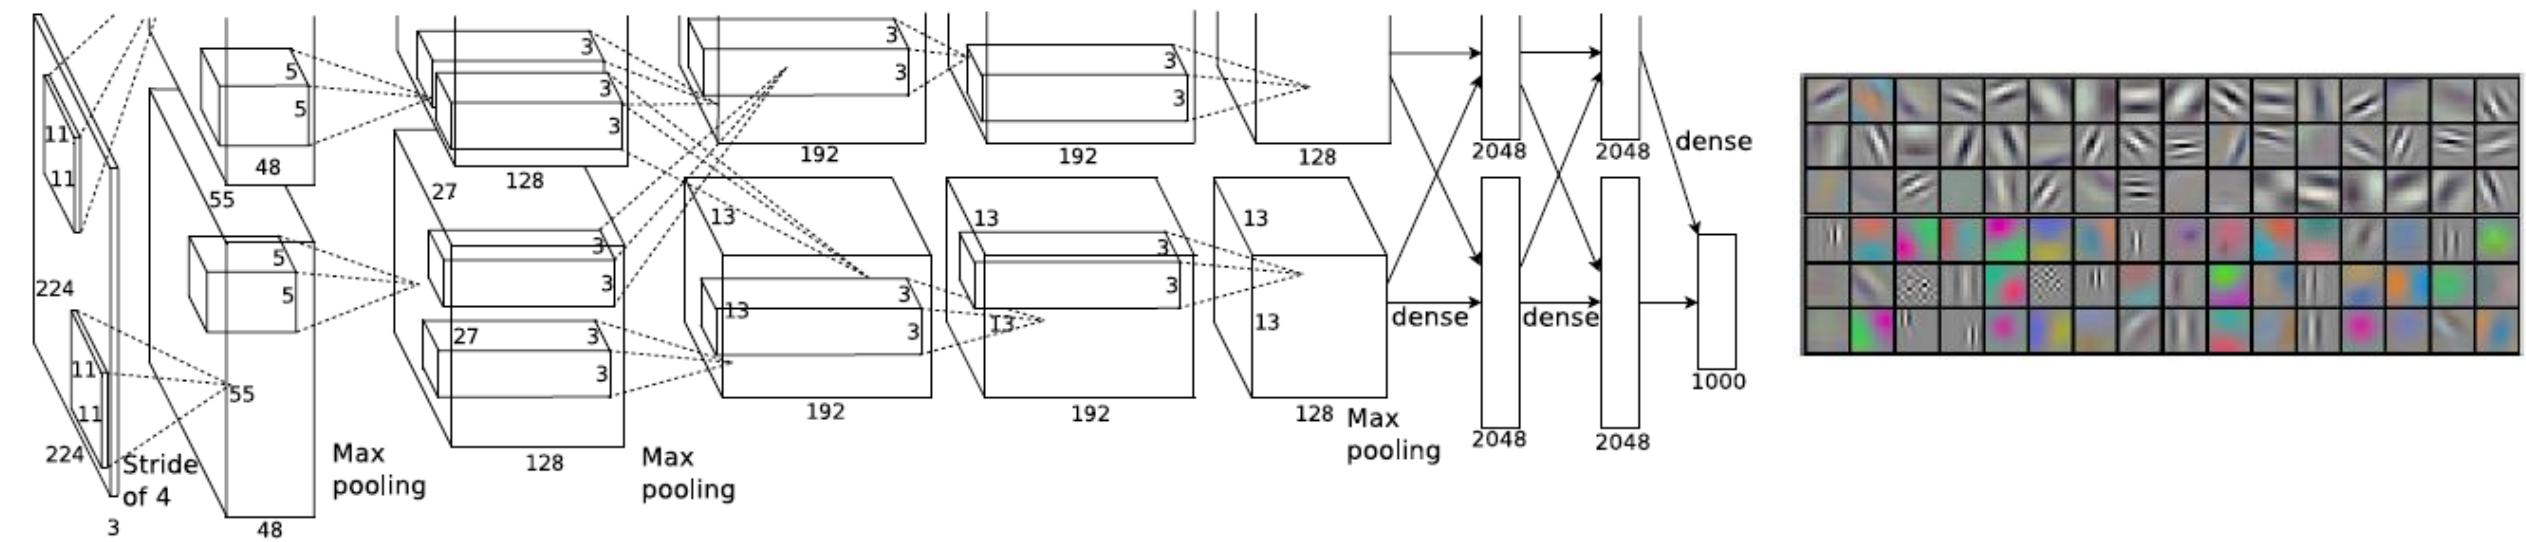
<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

# Popular CNN Models: LeNet



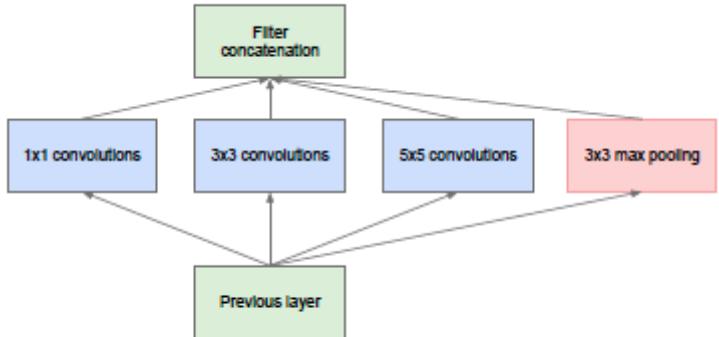
LeCun et al., Gradient-based learning applied to document recognition, 1998

# Popular CNN Models: AlexNet

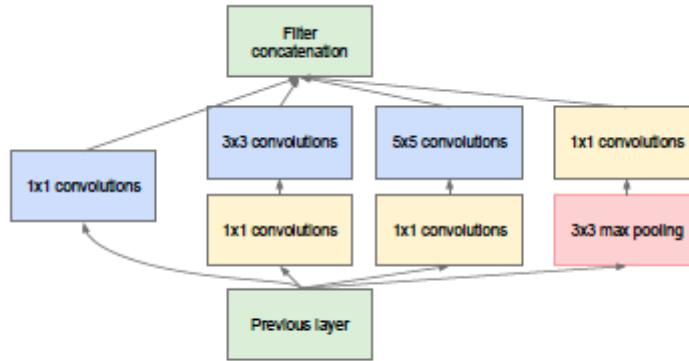


Krizhevsky et al., ImageNet Classification with Deep Convolutional Neural Networks, 2012

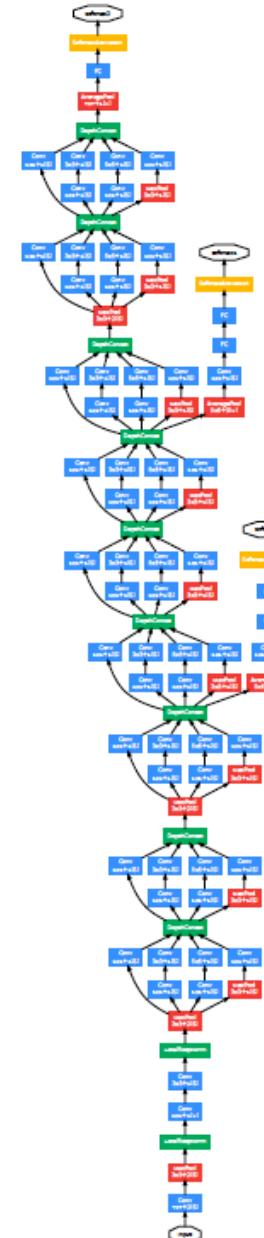
# Popular CNN Models: GoogLeNet (Inception)



(a) Inception module, naïve version

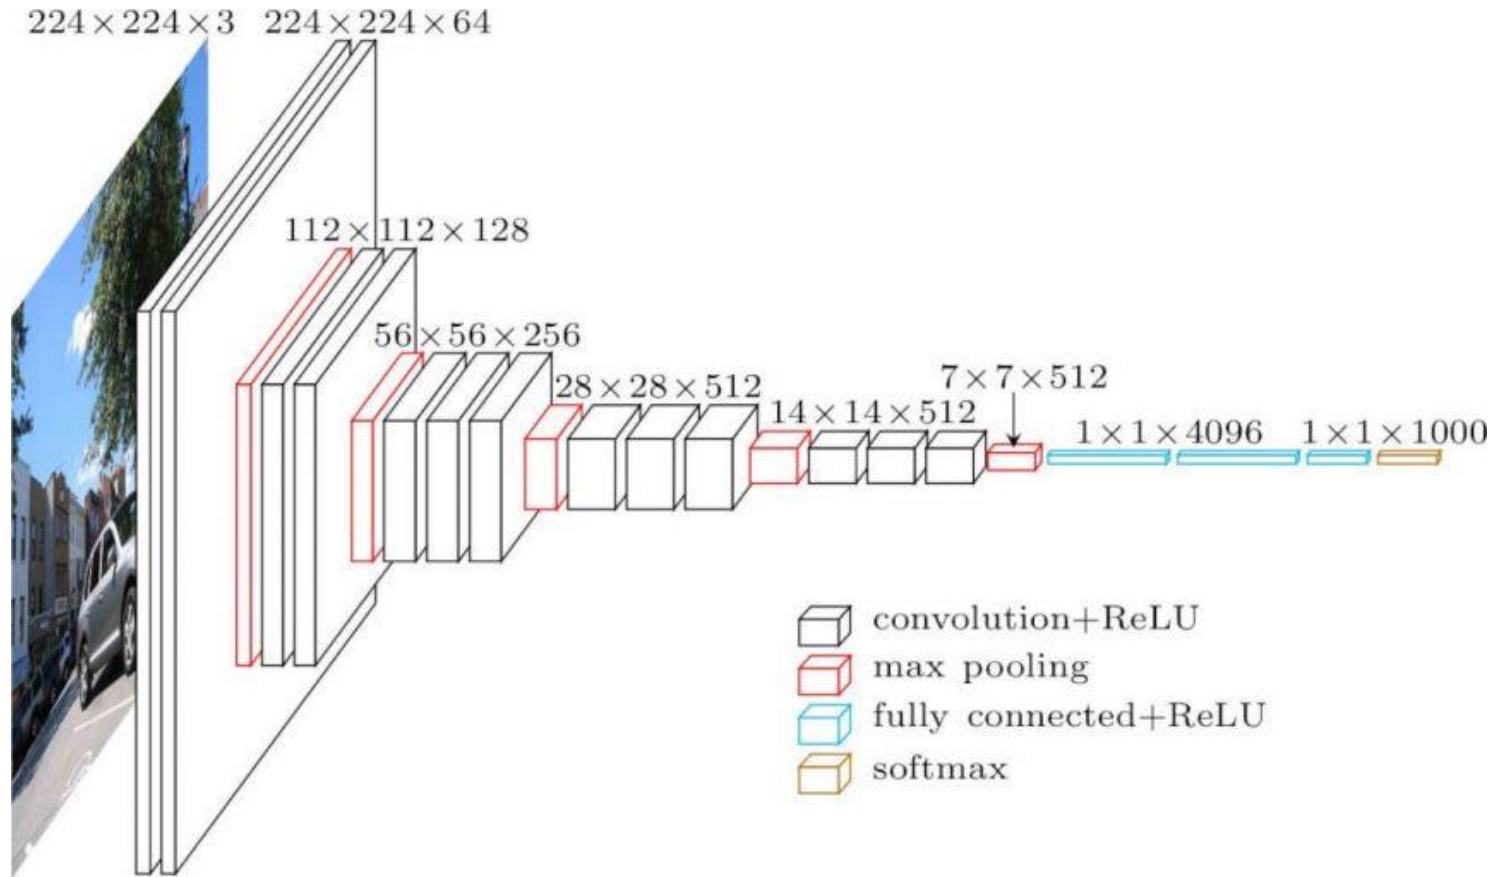


(b) Inception module with dimension reductions



Szegedy et al., Going deeper with convolutions, 2014

# Popular CNN Models: VGGNet



Simonyan & Zisserman, Very deep convolutional networks for large-scale image recognition, 2015  
Image from <https://neurohive.io/en/popular-networks/vgg16/>

# Popular CNN Models: ResNet

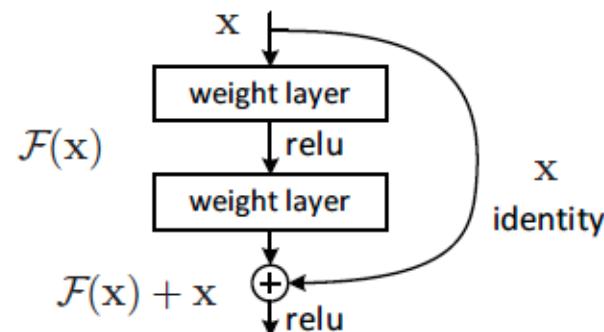
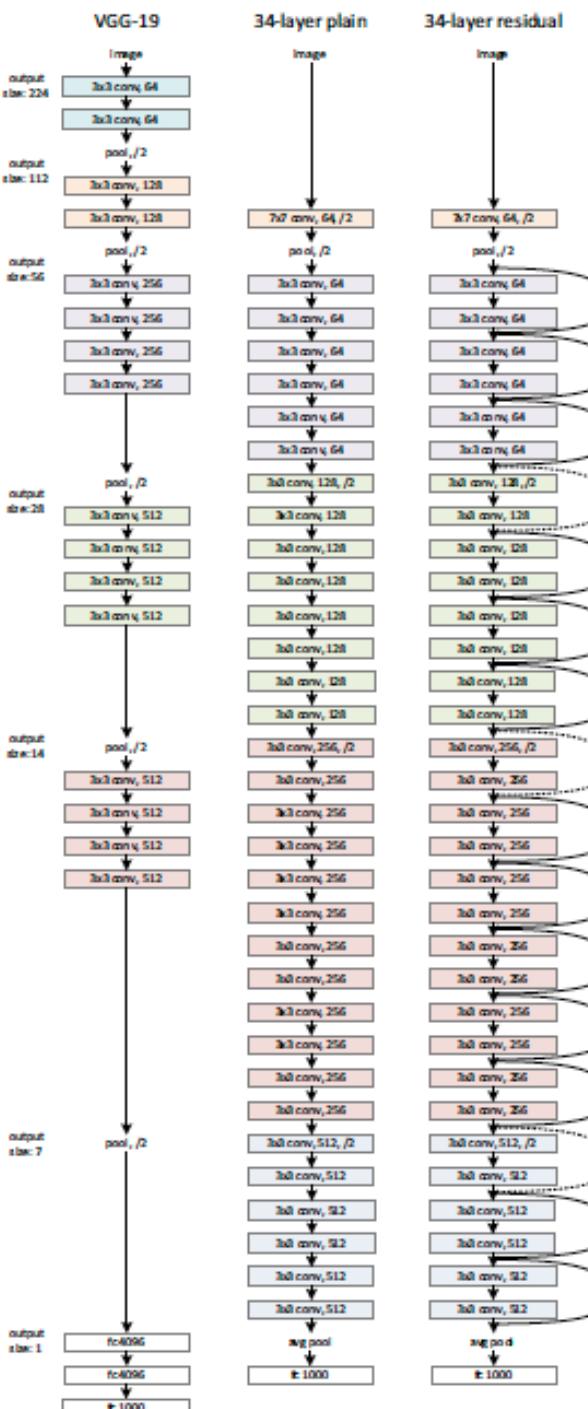
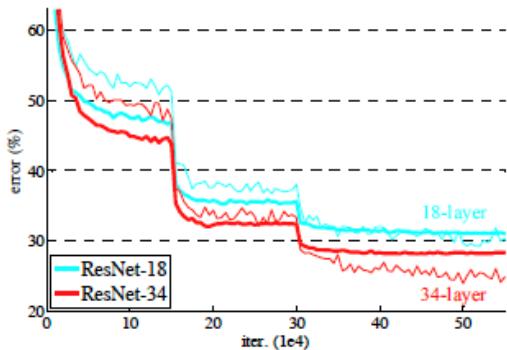
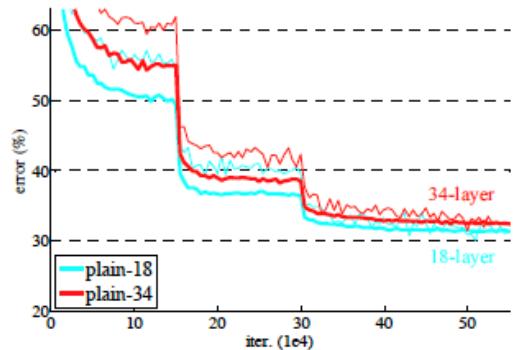


Figure 2. Residual learning: a building block.